

Konvolutivne Neuronske Mreže Convolutional Neural Networks, CNN

Predavač: Aleksandar Kovačević

Slajdovi preuzeti sa CS 231n, Stanford

<http://cs231n.stanford.edu/>

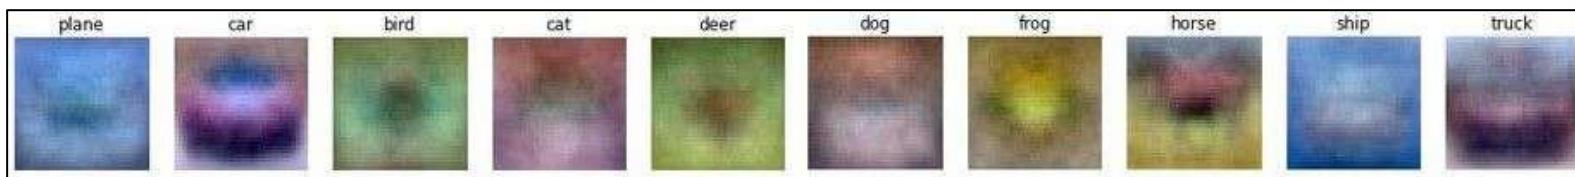
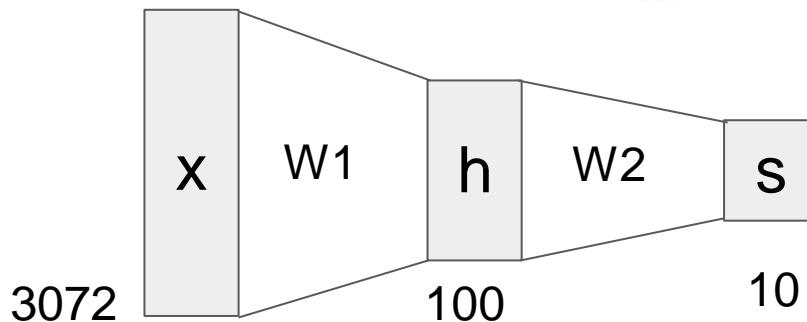
Do sada: Neuronske mreže

Linearna skor funkcija:

$$f = Wx$$

2-slojna Neuronska Mreža

$$f = W_2 \max(0, W_1 x)$$



Danas: Konvolutivne Neuronske Mreže

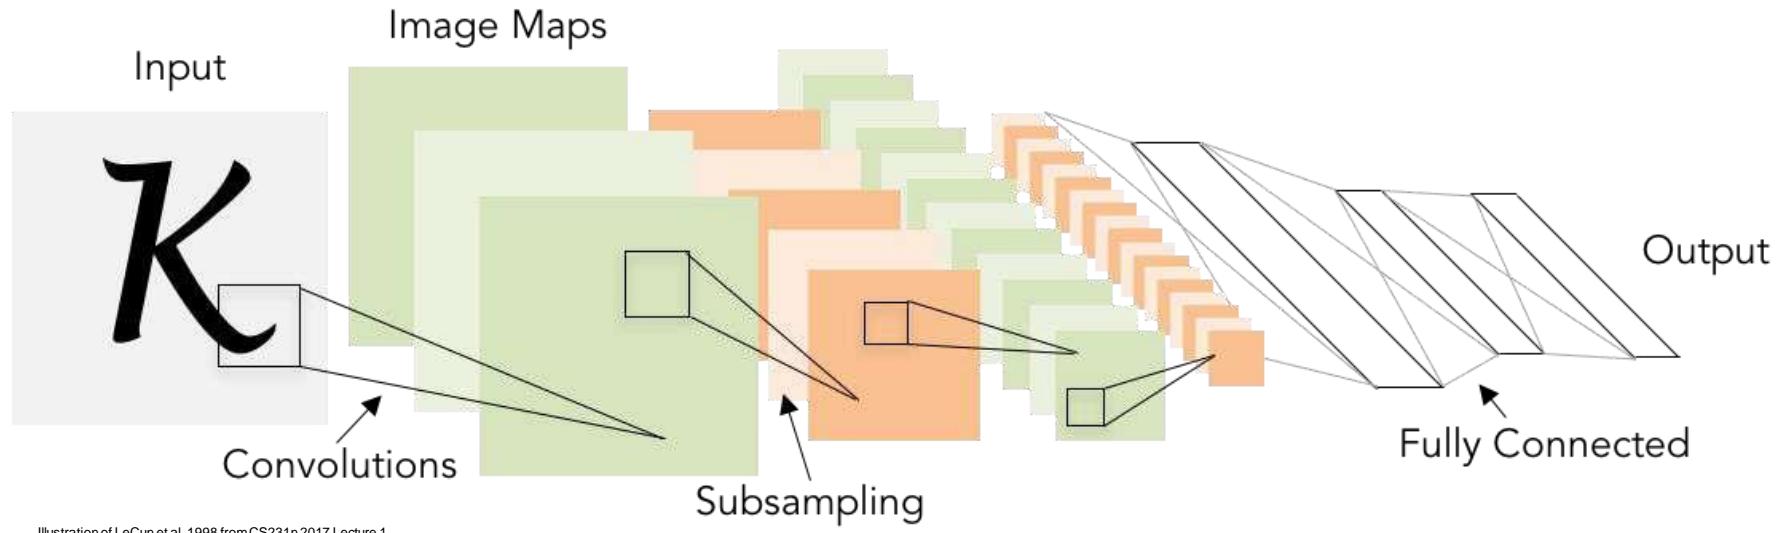


Illustration of LeCun et al. 1998 from CS231n 2017 Lecture 1

Malo istorije...

Mašina **Mark I Perceptron** je prva implementacija perceptron algoritma.

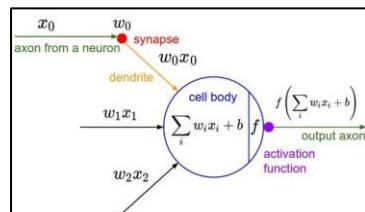
Mašina je bila povezana na kameru koja je koristila 20×20 kadmijum sulfid fotoćelije da proizvede sliku od 400 piksela.

mašina je mogla da prepoznaje slova abecede

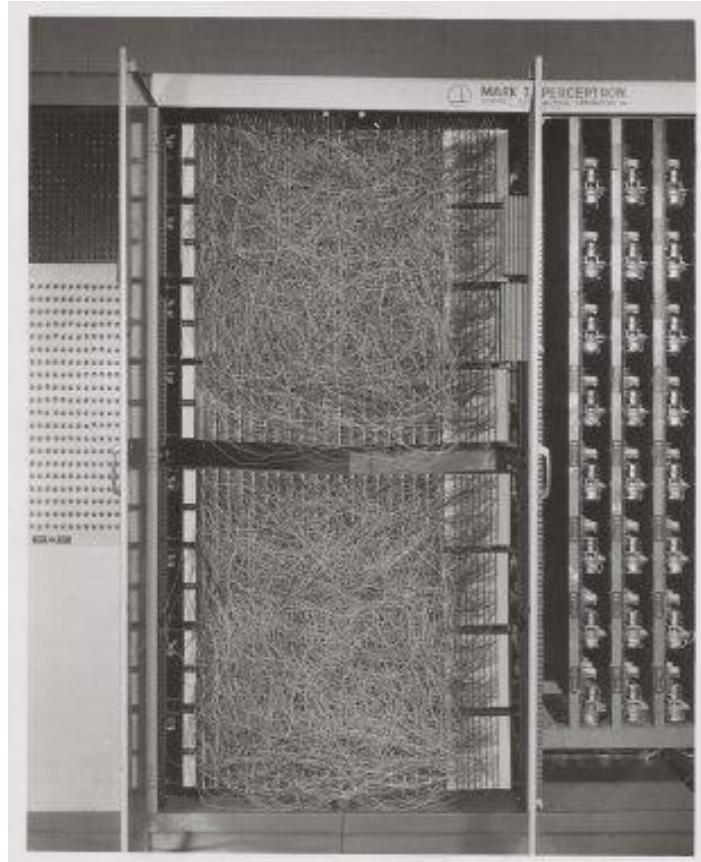
$$f(x) = \begin{cases} 1 & \text{if } w \cdot x + b > 0 \\ 0 & \text{otherwise} \end{cases}$$

pravilo za promenu težina:

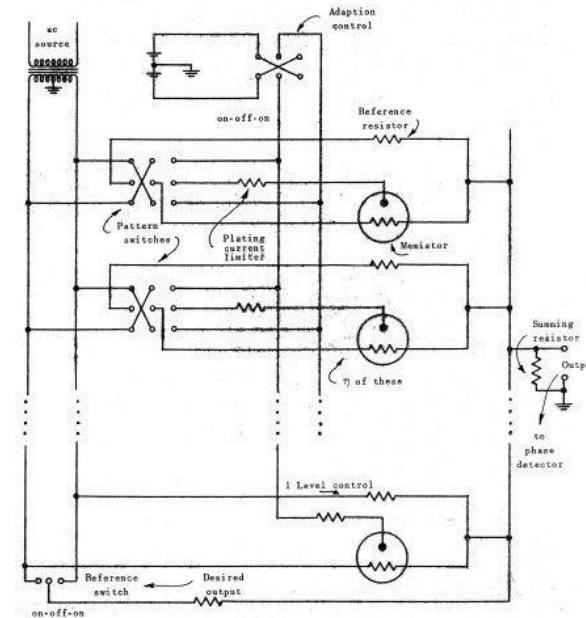
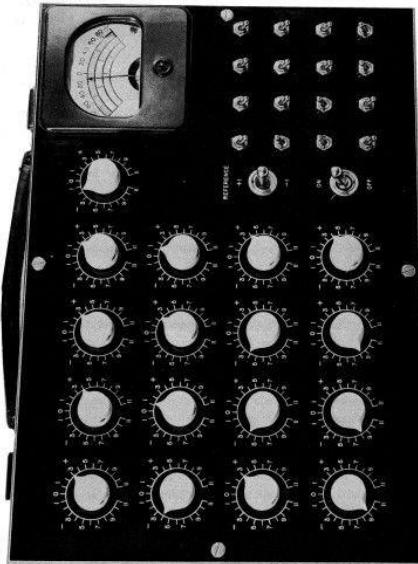
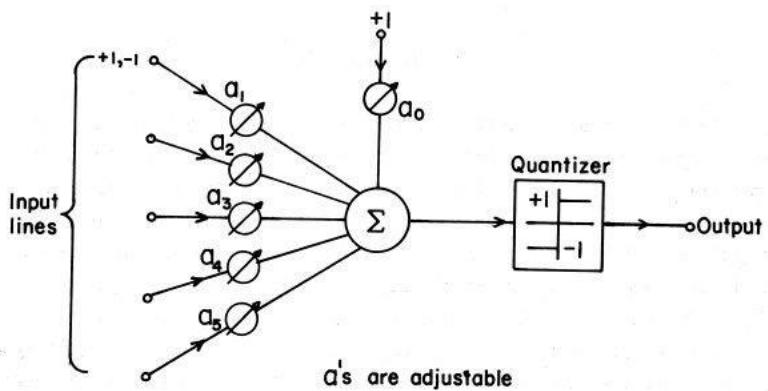
$$w_i(t+1) = w_i(t) + \alpha(d_j - y_j(t))x_{j,i}$$



Frank Rosenblatt, ~1957: Perceptron



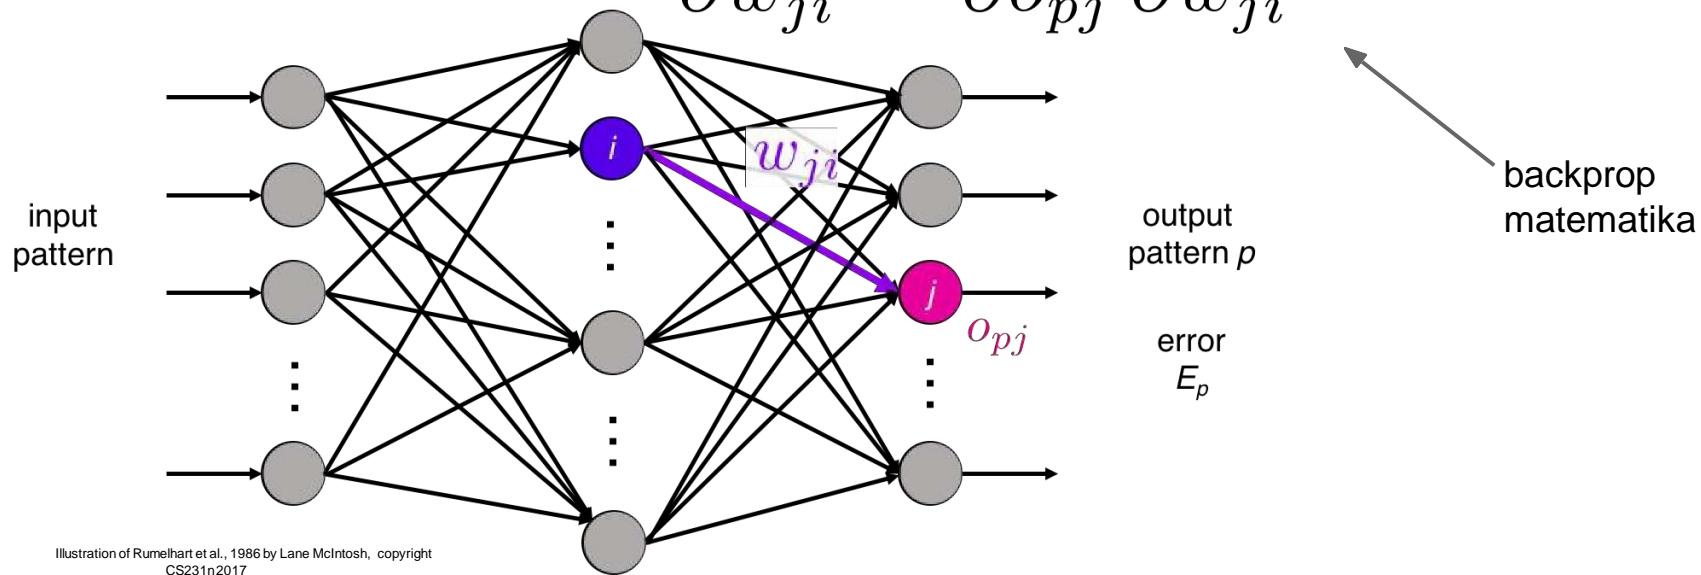
Malo istorije...



Widrow and Hoff, ~1960: Adaline/Madaline

Malo istorije...

$$\frac{\partial E_p}{\partial w_{ji}} = \frac{\partial E_p}{\partial o_{pj}} \frac{\partial o_{pj}}{\partial w_{ji}}$$



Rumelhart et al., 1986: Prvi put da je back-propagation postao popularan

Malo istorije...

[Hinton and Salakhutdinov 2006]

Počinje interesovanje za
Deep Learning

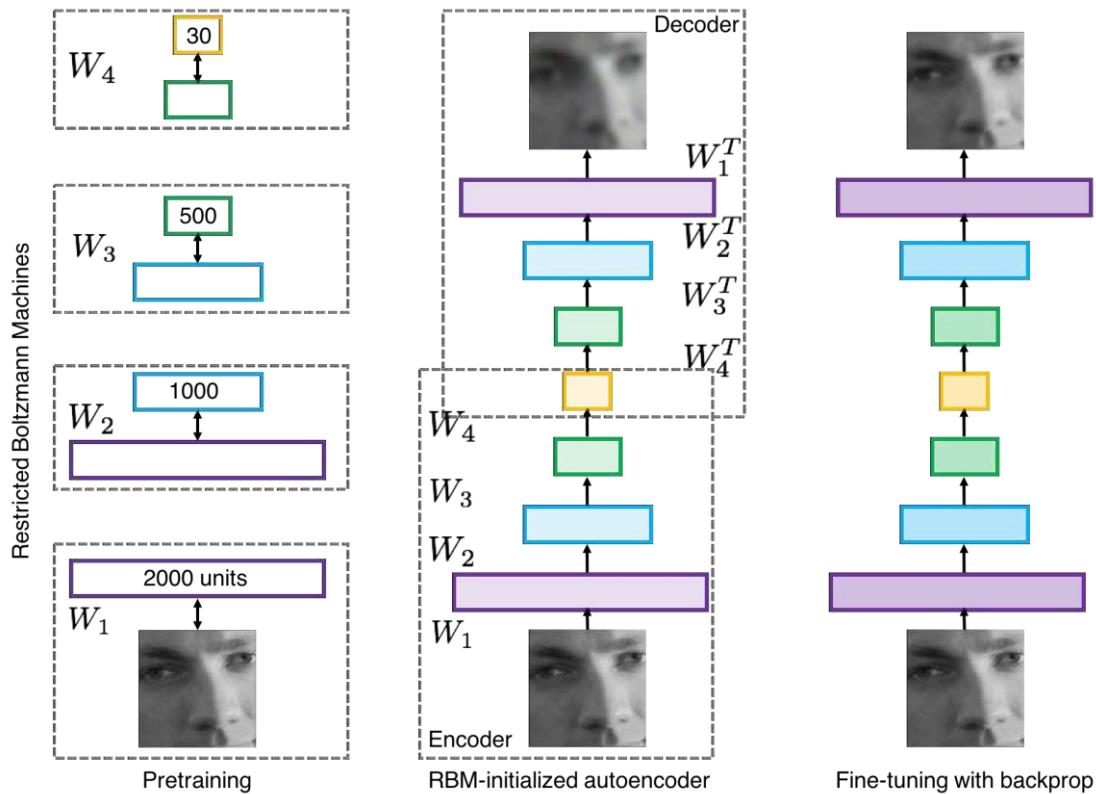


Illustration of Hinton and Salakhutdinov 2006 by Lane McIntosh,
copyright CS231n 2017

Prvi dobri rezultati

Acoustic Modeling using Deep Belief Networks

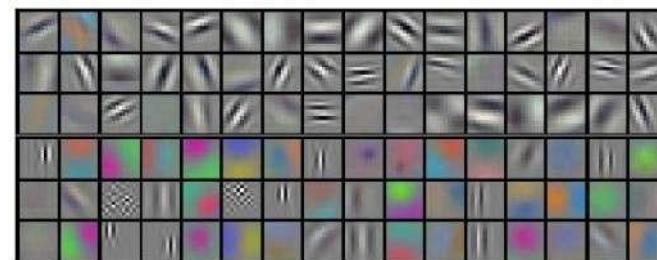
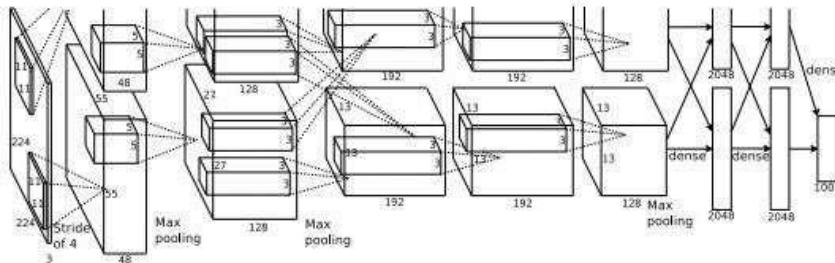
Abdel-rahman Mohamed, George Dahl, Geoffrey Hinton, 2010

Context-Dependent Pre-trained Deep Neural Networks for Large Vocabulary Speech Recognition

George Dahl, Dong Yu, Li Deng, Alex Acero, 2012

Imagenet classification with deep convolutional neural networks

Alex Krizhevsky, Ilya Sutskever, Geoffrey E Hinton, 2012



Figures copyright Alex Krizhevsky, Ilya Sutskever, and Geoffrey Hinton, 2012. Reproduced with permission.

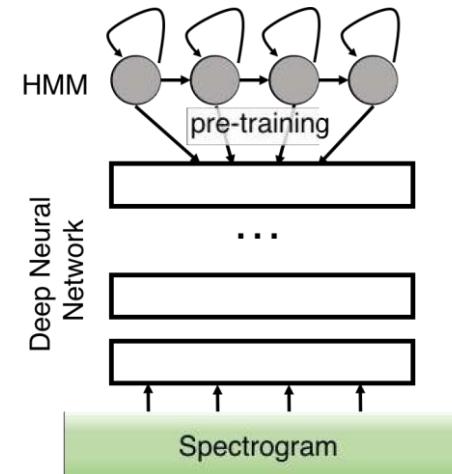


Illustration of Dahl et al. 2012 by Lane McIntosh, copyright CS231n 2017

Malo Istorije:

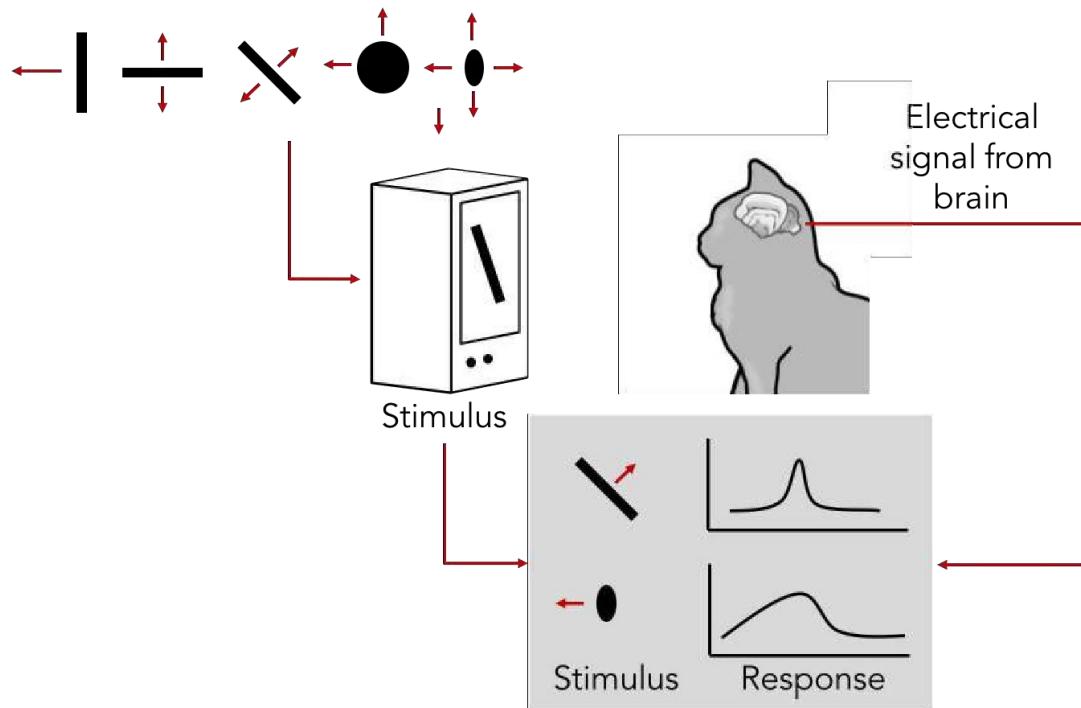
**Hubel & Wiesel,
1959**

RECEPTIVE FIELDS OF SINGLE
NEURONES IN
THE CAT'S STRIATE CORTEX

1962

RECEPTIVE FIELDS, BINOCULAR
INTERACTION
AND FUNCTIONAL ARCHITECTURE IN
THE CAT'S VISUAL CORTEX

1968...

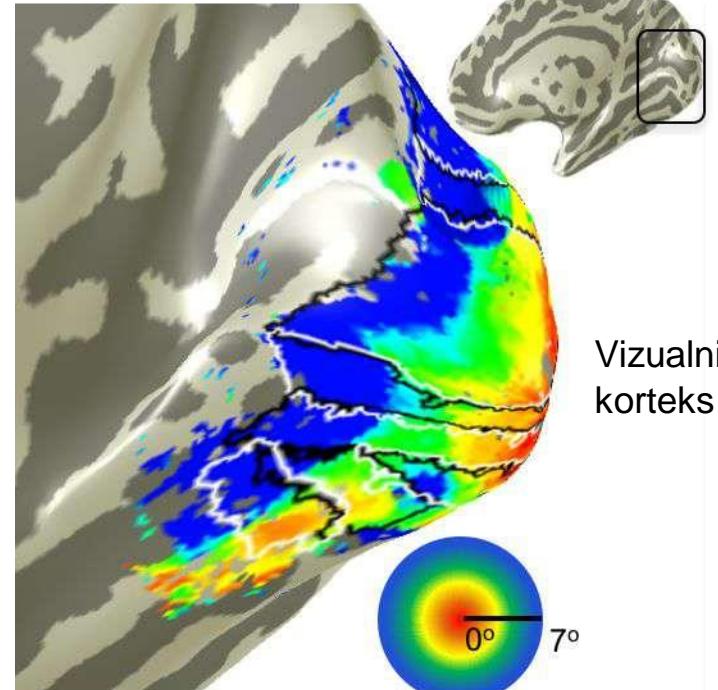
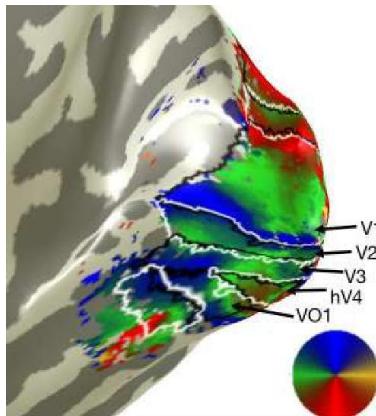


[Cat image](#) by CNX OpenStax is
licensed under CC BY 4.0; changes
made

Malo istorije...

Ljudski mozak

Topografska mapa korteksa:
ćelije koje su blizu u korteksu
procesiraju regije koji su
blizu u vizualnom polju



Hijerahjijska Organizacija

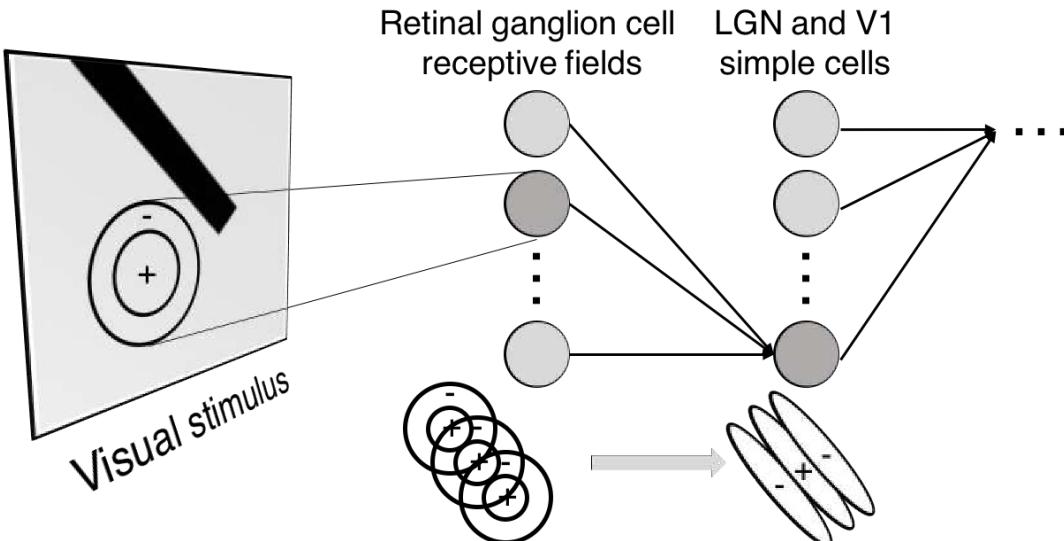
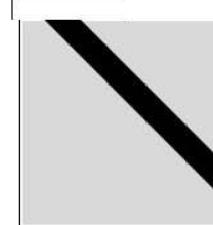


Illustration of hierarchical organization in early visual pathways
by Lane McIntosh, copyright CS231n 2017

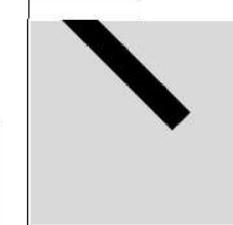
Simple cells:
Response to light orientation

Complex cells:
Response to light orientation and movement

Hypercomplex cells:
response to movement with an end point



No response



Response
(end point)

Malo istorije:

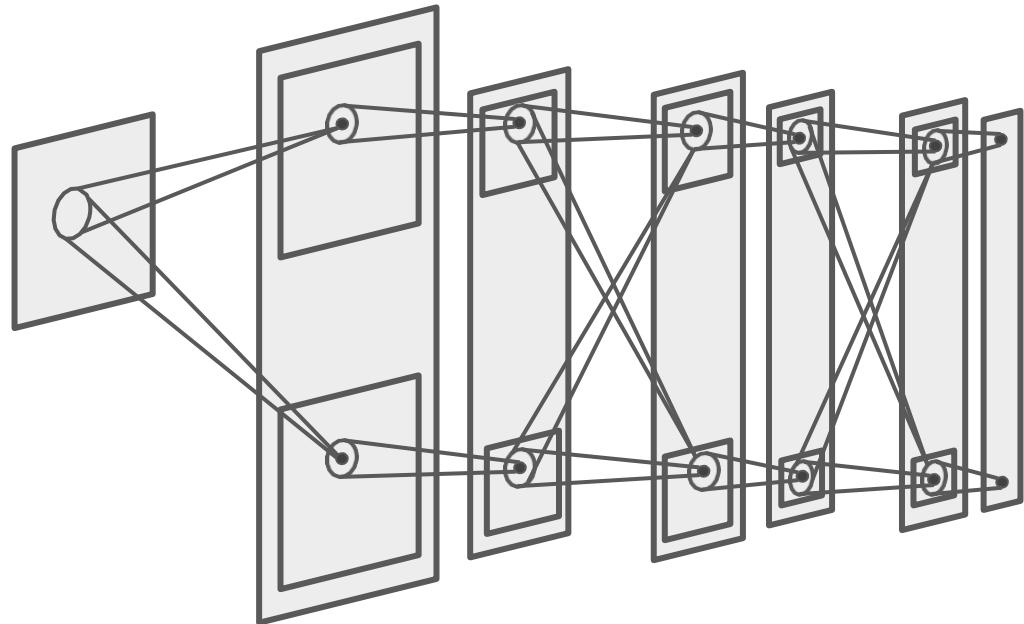
Neocognitron [Fukushima 1980]

“sendvič” arhitektura (SCSCSC...)

jednostavne ćelije: parametri koji se uče

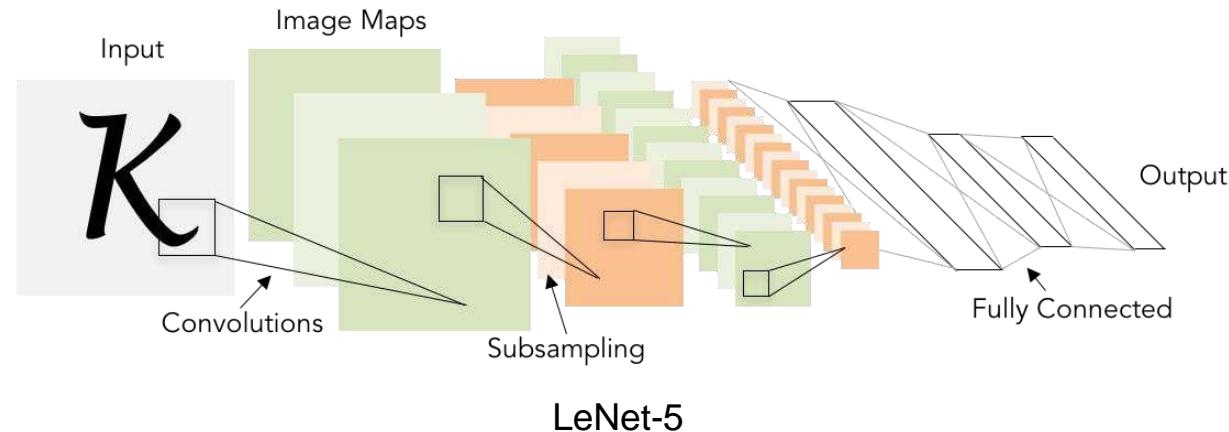
kompleksne ćelije: agregacija vrednosti

jednostavnih ćelija



Malo istorije: Gradient-based learning applied to document recognition

[LeCun, Bottou, Bengio, Haffner 1998]



Malo istorije: ImageNet Classification with Deep Convolutional Neural Networks [Krizhevsky, Sutskever, Hinton, 2012]

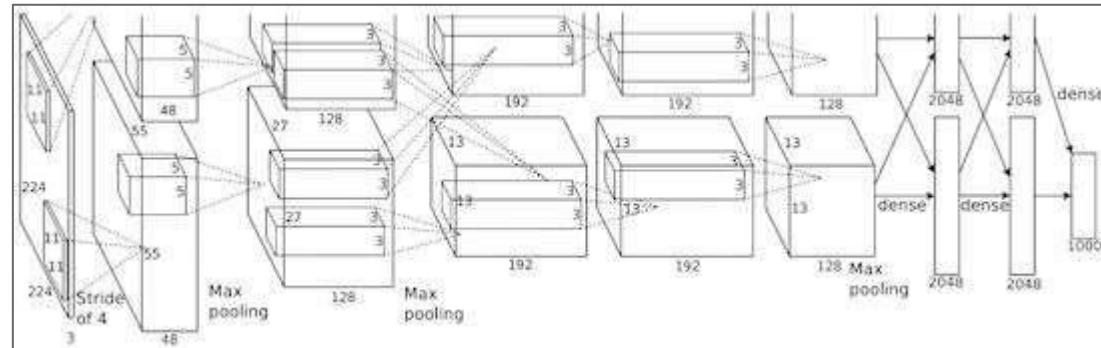
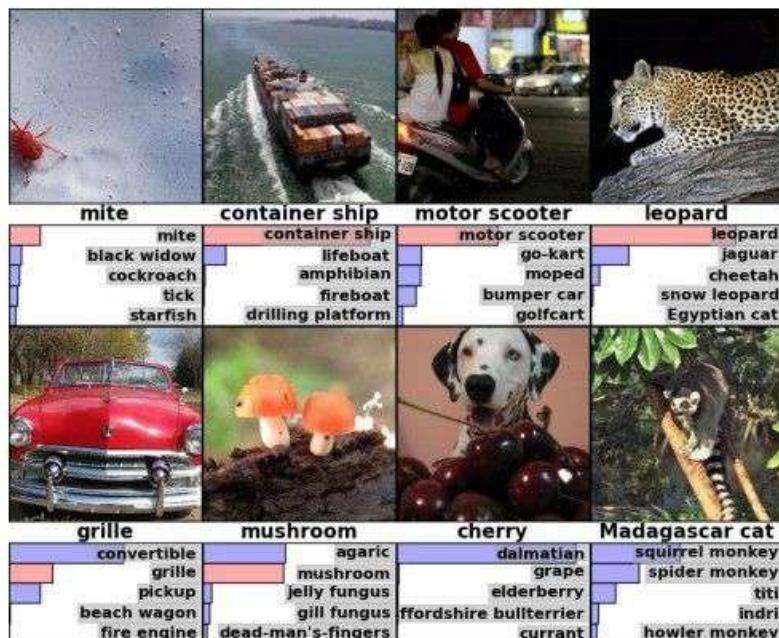


Figure copyright Alex Krizhevsky, Ilya Sutskever, and Geoffrey Hinton, 2012. Reproduced with permission.

“AlexNet”

Danas: Konvolutivne Mreže su svuda

Klasifikacija



Pronalaženje



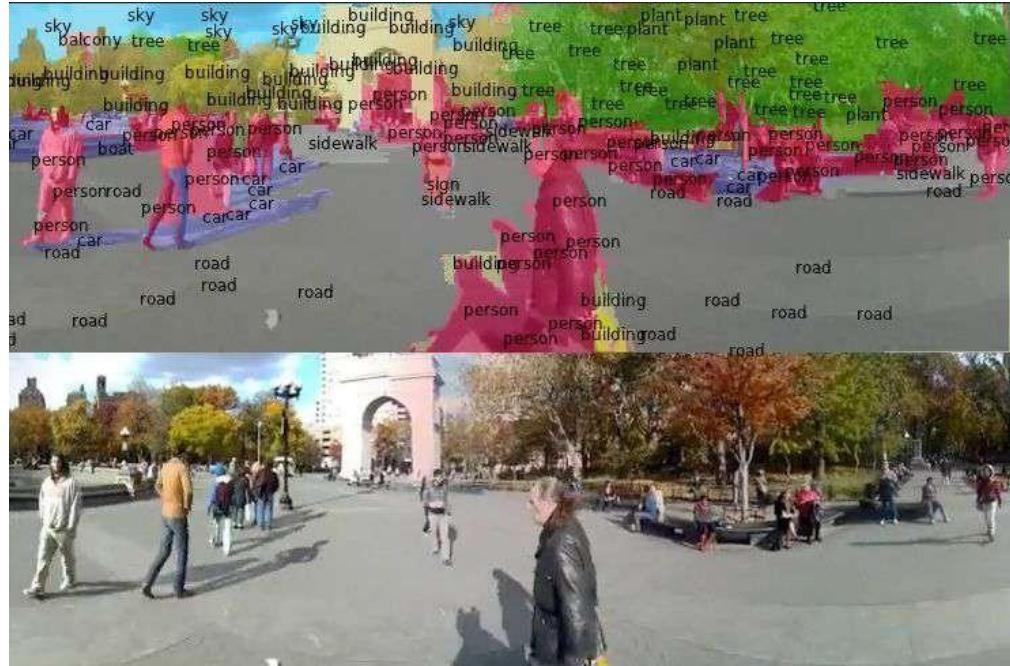
Figures copyright Alex Krizhevsky, Ilya Sutskever, and Geoffrey Hinton, 2012. Reproduced with permission.

Danas: Konvolutivne Mreže su svuda

Detekcija



Segmentacija



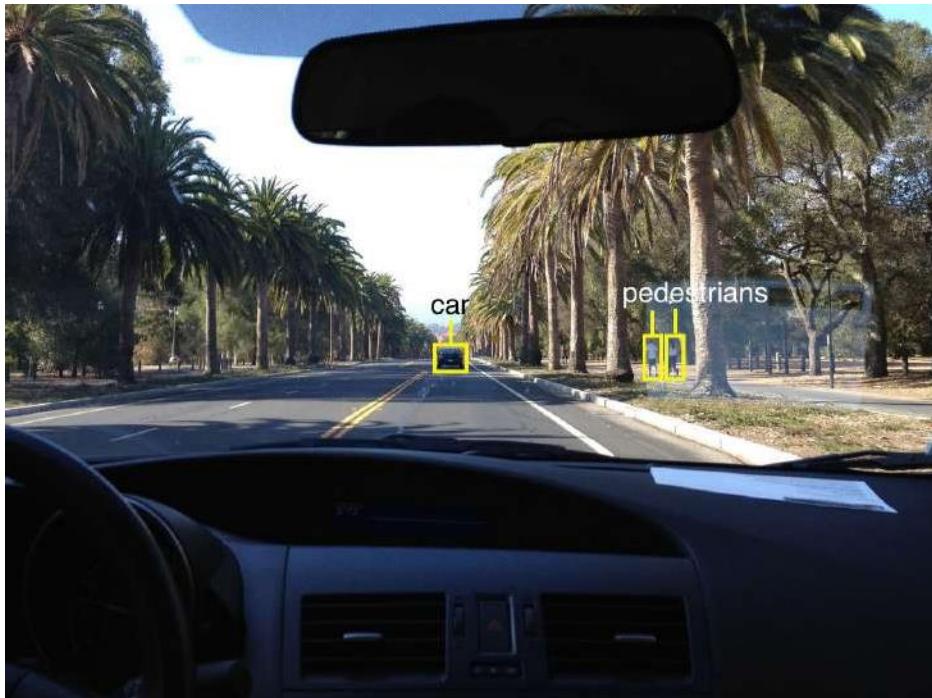
Figures copyright Shaoqing Ren, Kaiming He, Ross Girshick, Jian Sun, 2015.
Reproduced with permission.

[Faster R-CNN: Ren, He, Girshick, Sun 2015]

Figures copyright Clement Farabet,
2012. Reproduced with permission.

[Farabet et al., 2012]

Danas: Konvolutivne Mreže su svuda



self-driving cars

Photo by Lane McIntosh. Copyright CS231n
2017.

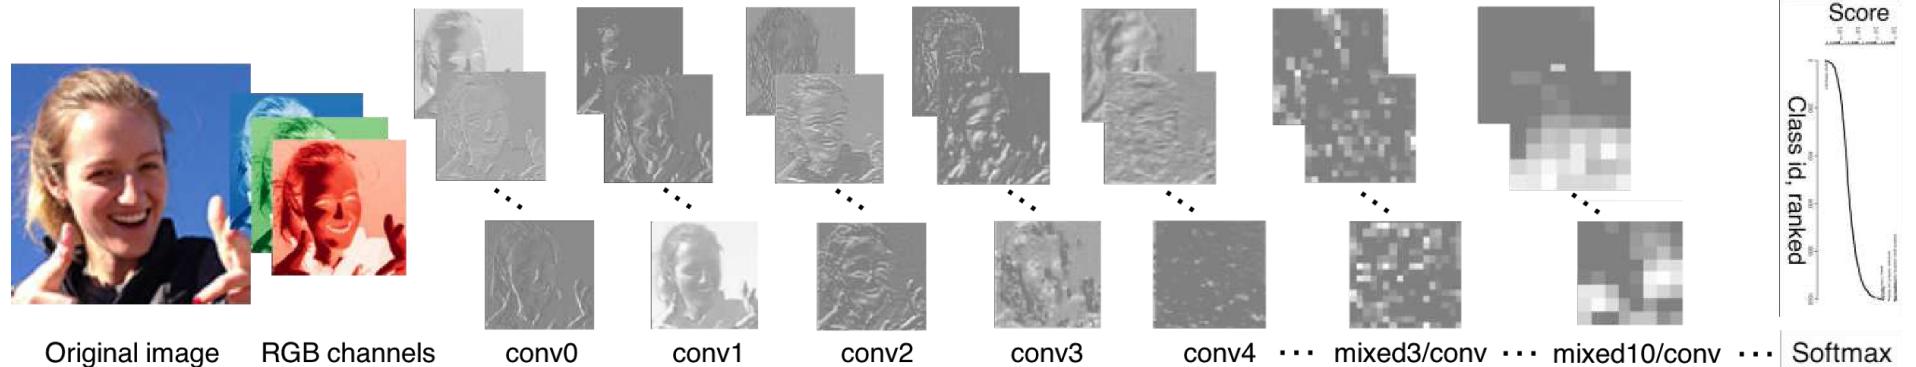


[This image](#) by
GBPublic_PR is licensed
under [CC-BY 2.0](#)

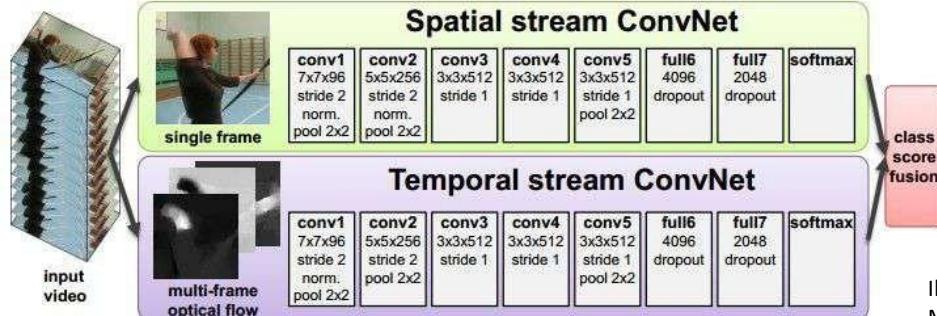
NVIDIA Tesla line

Za *embedded sisteme* tipično se koriste NVIDIA Tegra kartice, sa integriranim GPU i procesoru zasnovanom na ARM arhitekturi.

Danas: Konvolutivne Mreže su svuda



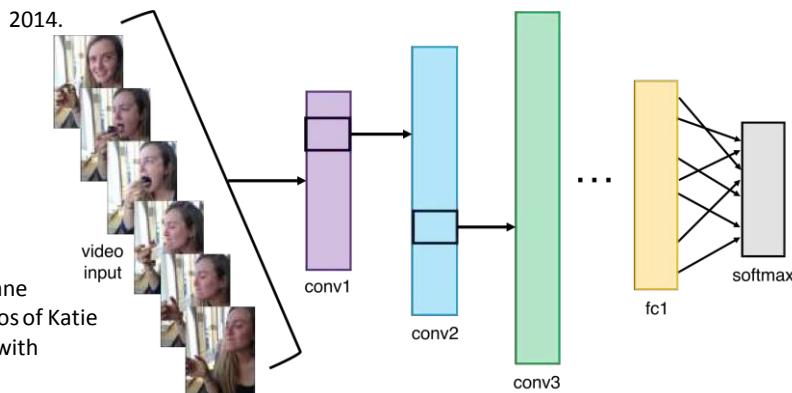
Activations of [inception-v3 architecture](#) [Szegedy et al. 2015] to image of Emma McIntosh, used with permission. Figure and architecture not from Taigman et al.



[Simonyan et al. 2014]

Figures copyright Simonyan et al., 2014. Reproduced with permission.

Illustration by Lane McIntosh, photos of Katie Cumnock used with permission.

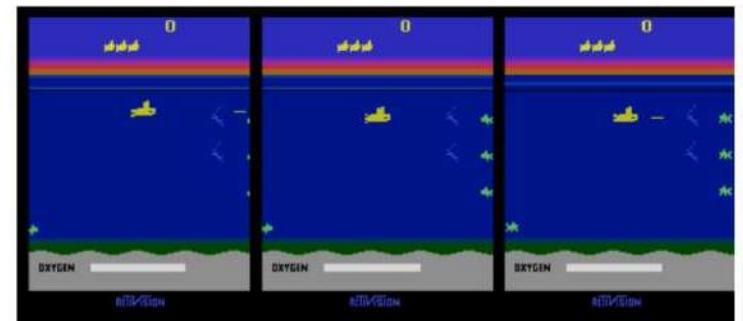
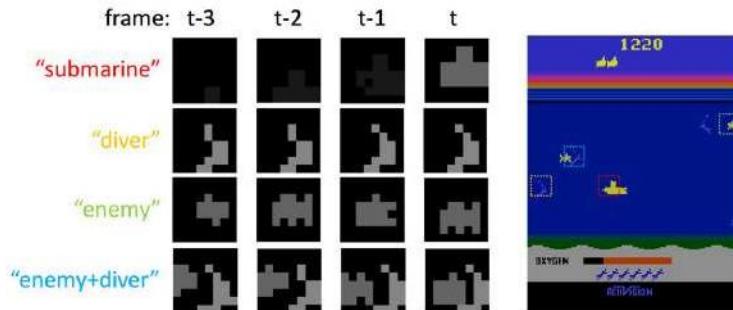


Danas: Konvolutivne Mreže su svuda



[Toshev, Szegedy 2014]

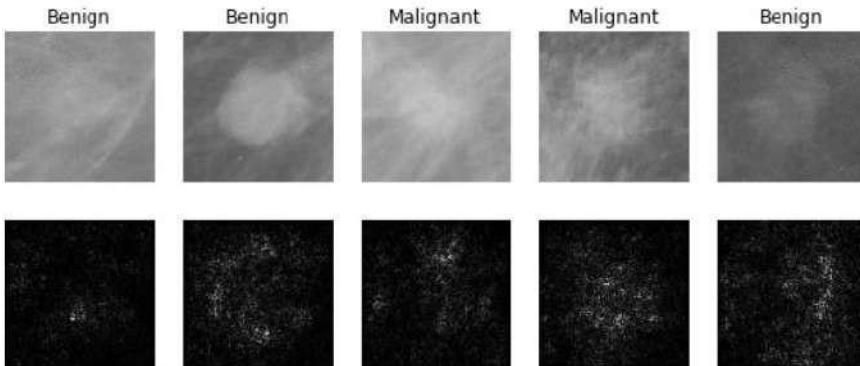
Images are examples of pose estimation, not actually from Toshev & Szegedy 2014. Copyright Lane McIntosh.



[Guo et al. 2014]

Figures copyright Xiaoxiao Guo, Satinder Singh, Honglak Lee, Richard Lewis, and Xiaoshi Wang, 2014. Reproduced with permission.

Danas: Konvolutivne Mreže su svuda



[Levy et al. 2016]

Figure copyright Levy et al. 2016.
Reproduced with permission.



[Dieleman et al. 2014]

From left to right: [public domain by NASA](#), usage [permitted](#) by ESA/Hubble,
[public domain by NASA](#), and [public domain](#).



Photos by Lane McIntosh.
Copyright CS231n 2017.

[Sermanet et al. 2011]
[Ciresan et al.]

[This image](#) by Christin Khan is in the public domain and originally came from the U.S. NOAA.



Whale recognition, Kaggle Challenge

Photo and figure by Lane McIntosh; not actual example from Mnih and Hinton, 2010 paper.



Mnih and Hinton, 2010

Bez grešaka



A white teddy bear sitting in the grass

Manje greške



A man in a baseball uniform throwing a ball

Ima neke veze sa slikom



A woman is holding a cat in her hand

Image Captioning

[Vinyals et al., 2015]
[Karpathy and Fei-Fei, 2015]



A man riding a wave on top of a surfboard



A cat sitting on a suitcase on the floor



A woman standing on a beach holding a surfboard

All images are CC0 Public domain:
<https://pixabay.com/en/luggage-antique-cat-1643010/>
<https://pixabay.com/en/teddy-plush-bears-cute-teddy-bear-1623436/>
<https://pixabay.com/en/surf-wave-summer-sport-litoral-1668716/>
<https://pixabay.com/en/woman-female-model-portrait-adult-983967/>
<https://pixabay.com/en/handstand-lake-meditation-496008/>
<https://pixabay.com/en/baseball-player-shortstop-infield-1045263/>

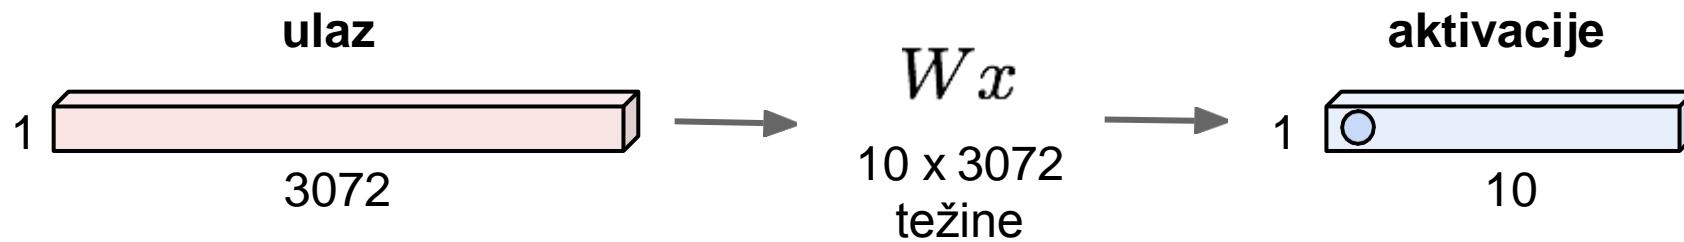
Captions generated by Justin Johnson using [Neuraltalk2](#)

Konvolutivne Neuronske Mreže

(Prvo bez analogije sa ljudskim mozgom)

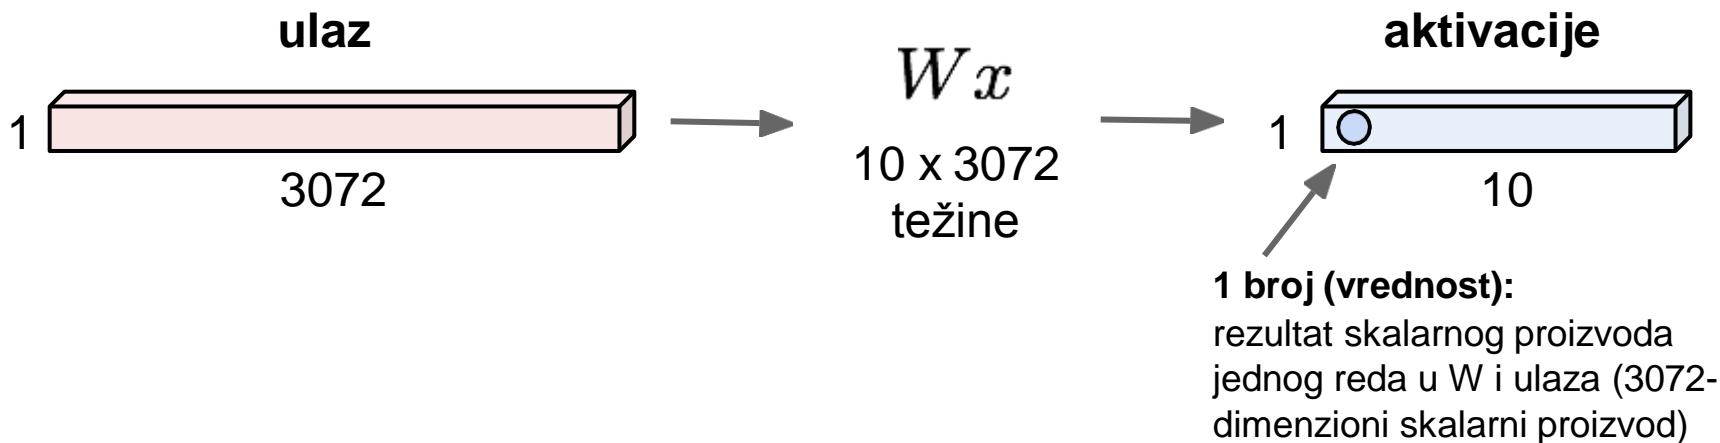
Potpuno Povezan Sloj

32x32x3 slika -> razvučemo u niz 3072×1



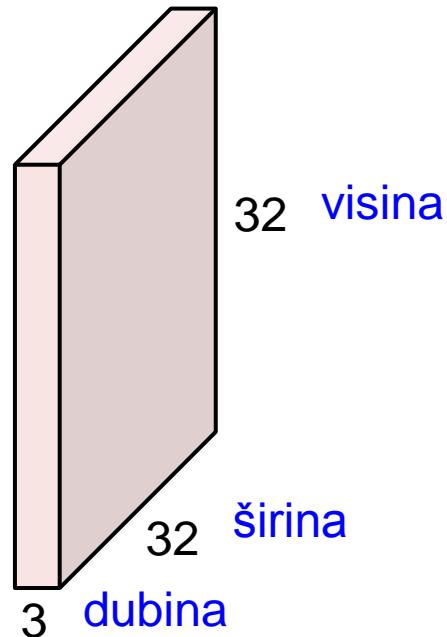
Potpuno Povezan Sloj

32x32x3 slika -> razvučemo u niz 3072×1



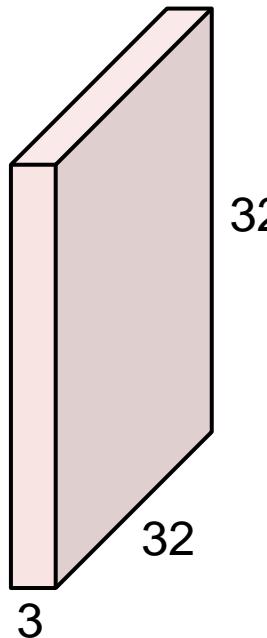
Konvolutivn Sloj

32x32x3 slika -> očuvavamo prostornu strukturu (slika ostaje 2d, a ne 1d niz)



Konvolutivni Sloj

32x32x3 slika

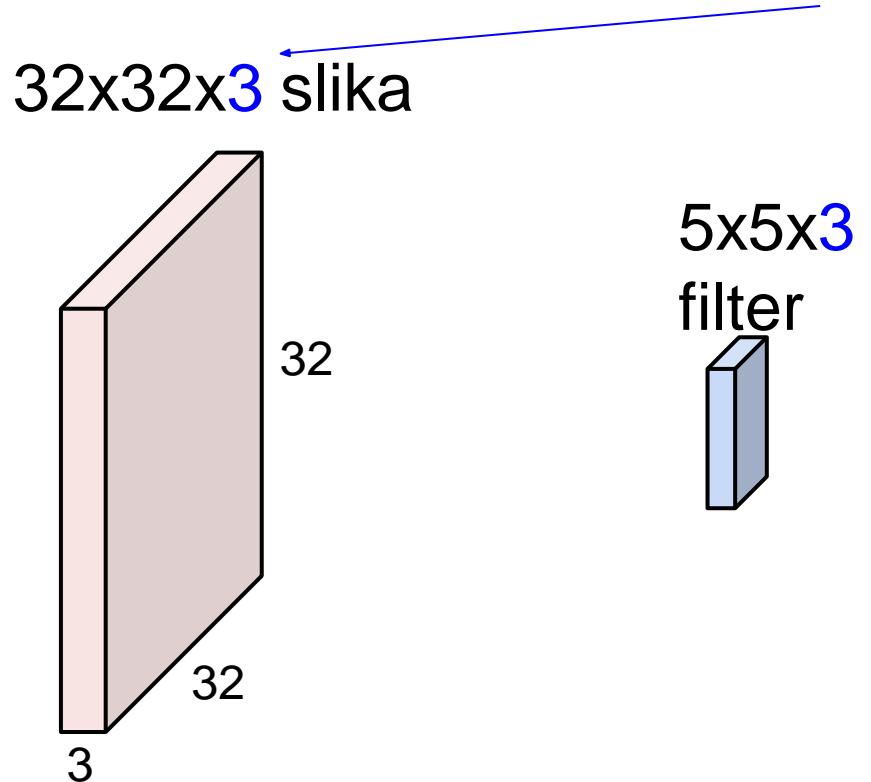


5x5x3
filter



Vršimo Konvoluciju filtera sa slikom
tj. "vučemo" filter prostorno po slici i
računamo skalarni proizvod

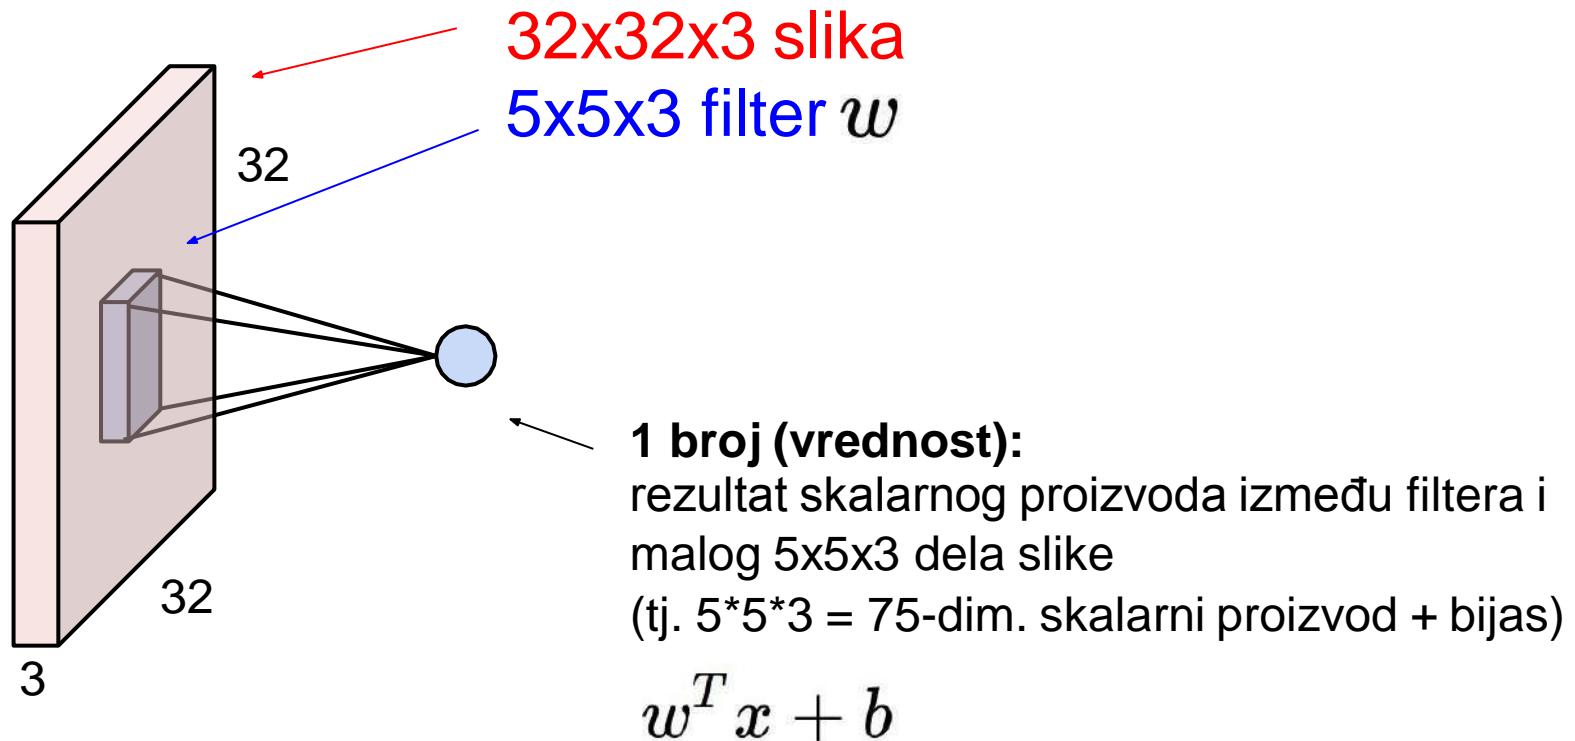
Konvolutivni Sloj



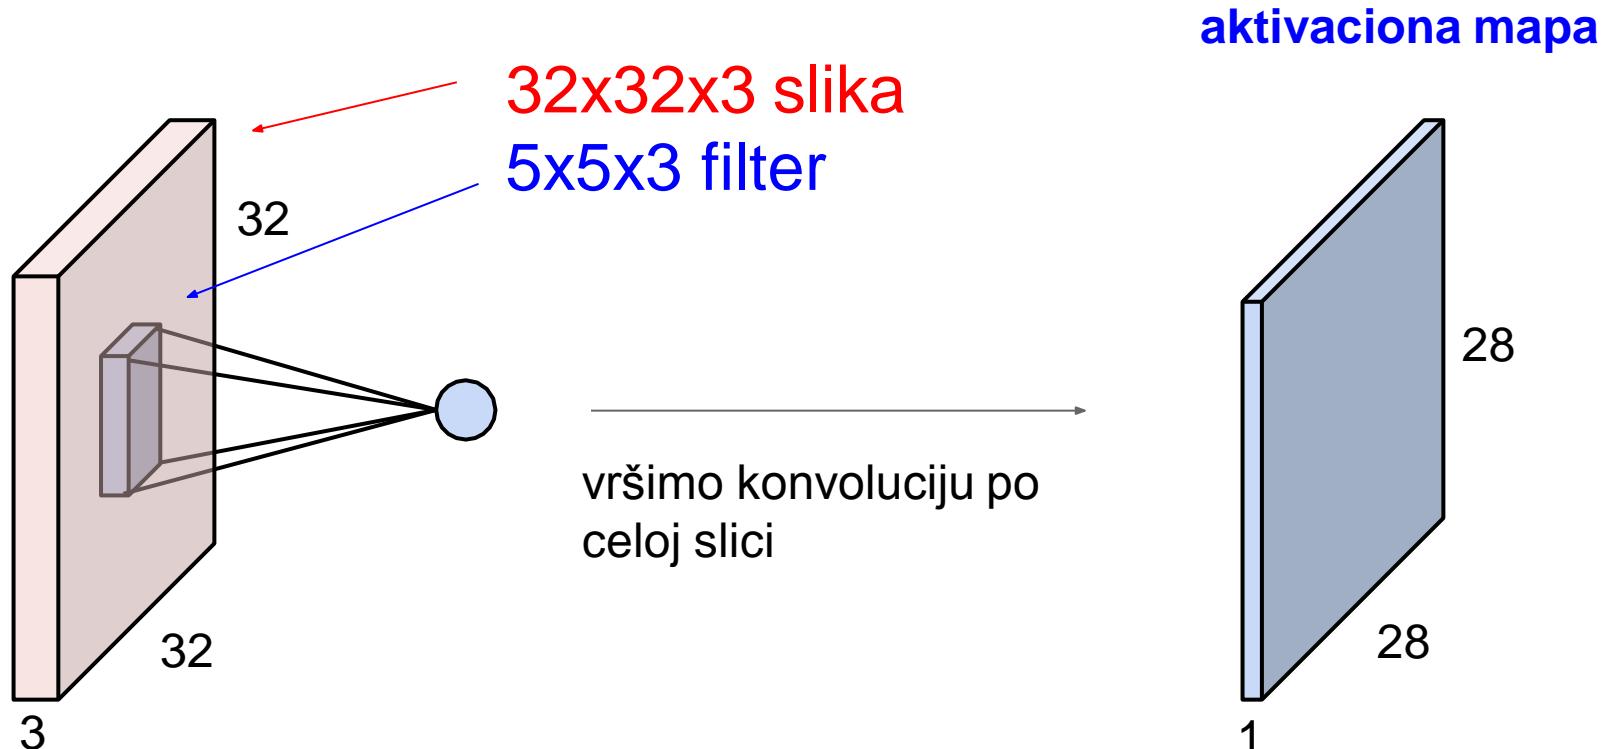
Filteri su uvek iste dubine
kao i slika

Vršimo Konvoluciju filtera sa slikom
tj. "vučemo" filter prostorno po slici i
računamo skalarni proizvod

Konvolutivni Sloj

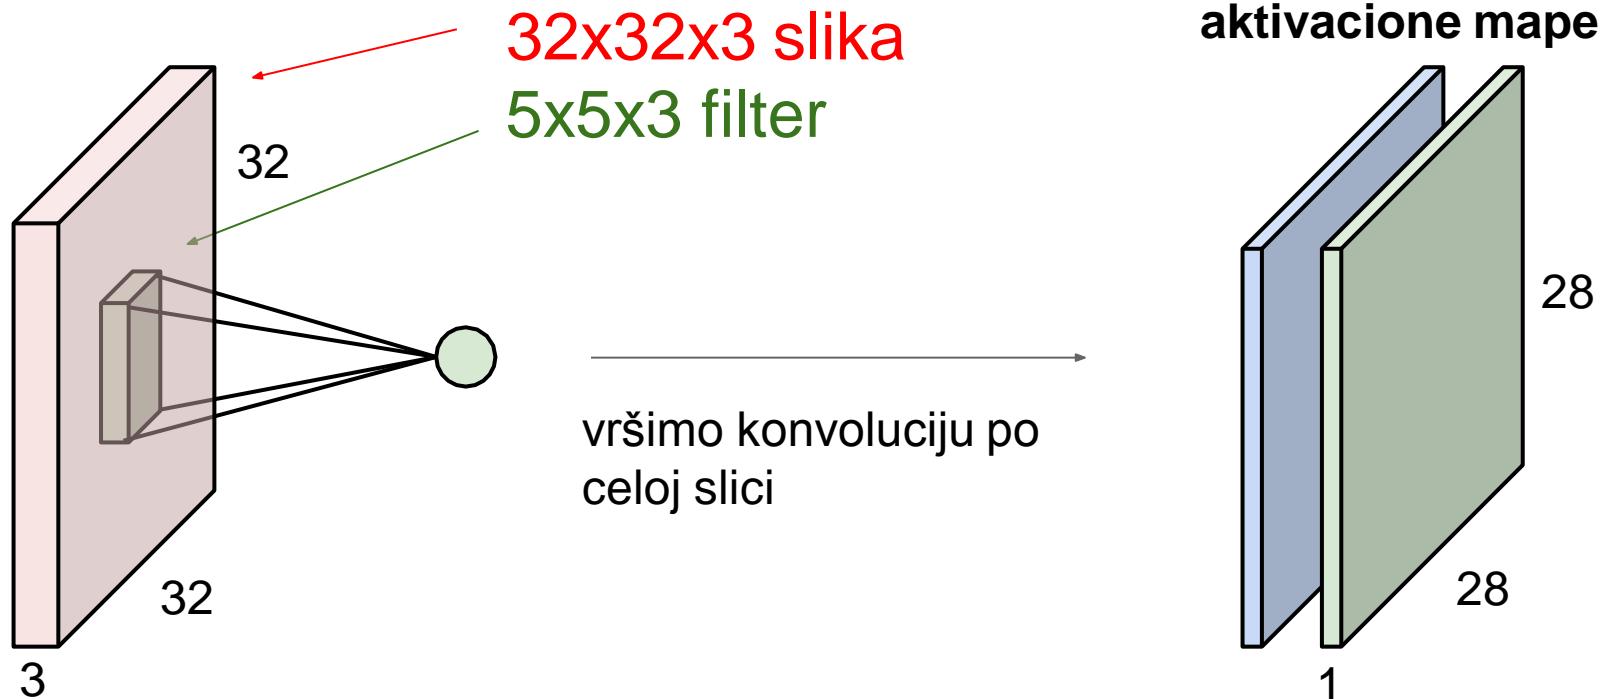


Konvolutivni Sloj

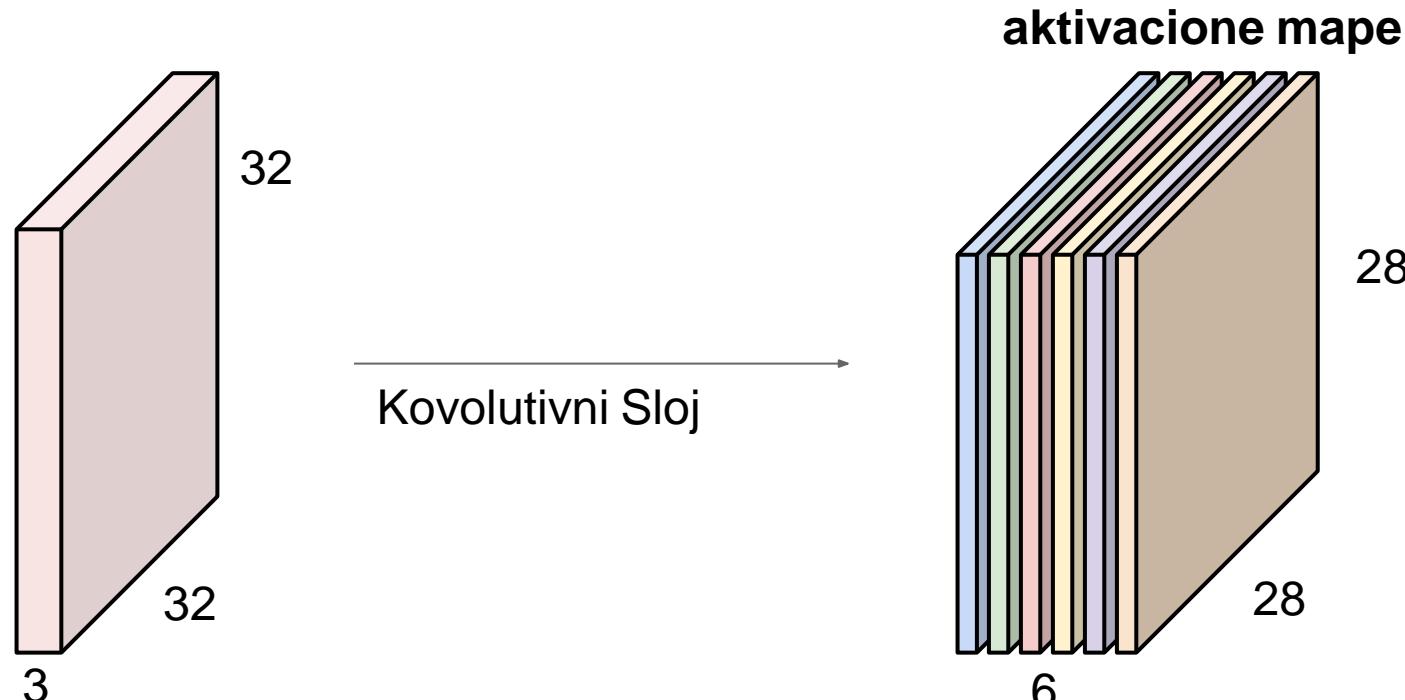


Konvolutivni Sloj

recimo da imamo, zeleni
filter



Na primer, ako imamo 6 5x5 filtera, dobićemo 6 odvojenih aktivacionih mapa:



Slažemo aktivacione mape jednu na drugu i tako dobijamo "novu sliku" veličine 28x28x6!

Primer konvolucije

Narednih nekoliko slajdova zasnovano je na

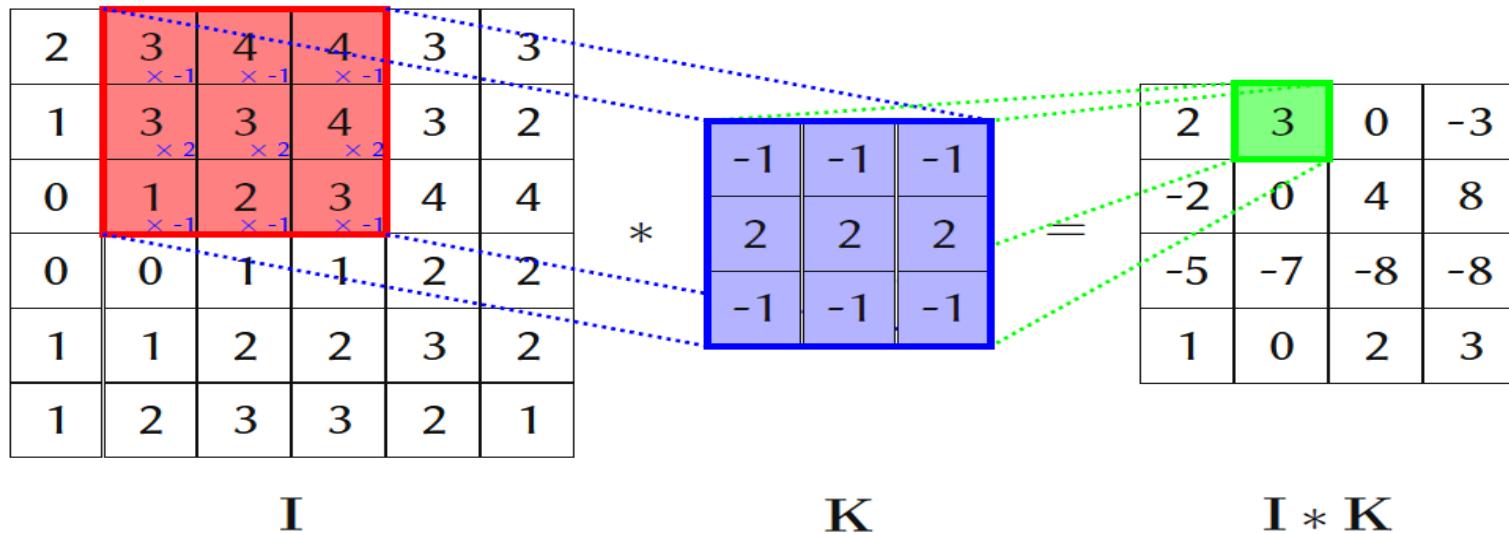
http://www.csnedelja.mg.edu.rs/static/resources/v3.0/sre1_neuralne_pv.pdf

Primer konvolucije

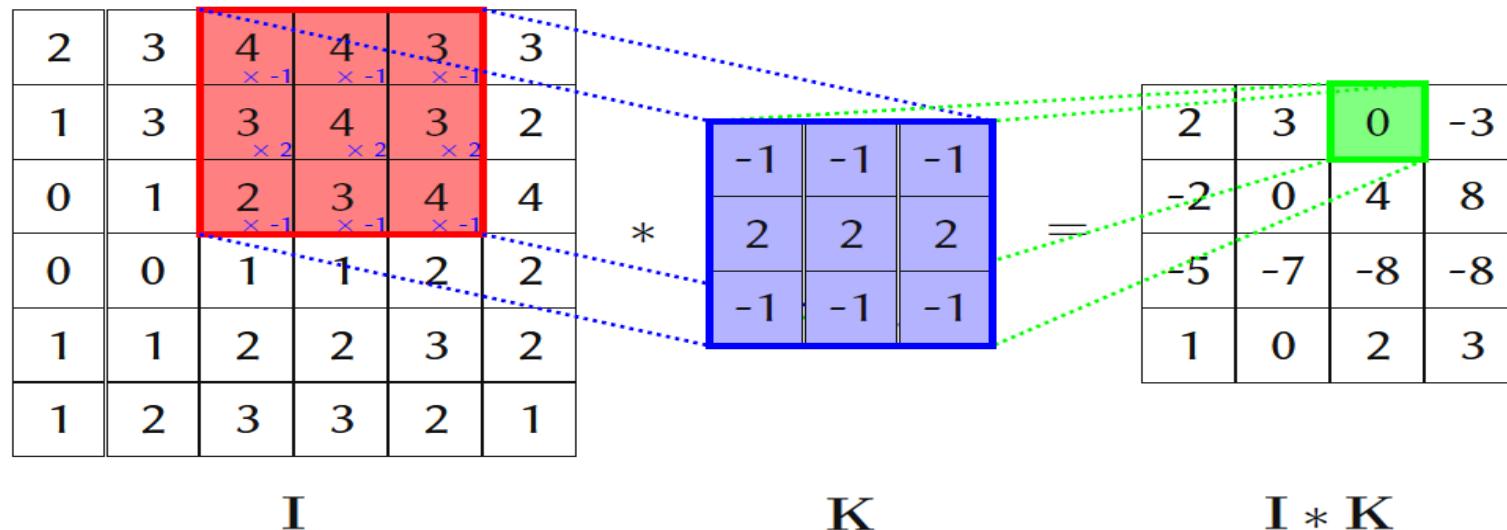
$$\begin{array}{c}
 \begin{array}{|c|c|c|c|c|c|} \hline
 2 & 3 & 4 & 4 & 3 & 3 \\ \hline
 \times -1 & \times -1 & \times -1 & & & \\ \hline
 1 & 3 & 3 & 4 & 3 & 2 \\ \hline
 \times 2 & \times 2 & \times 2 & & & \\ \hline
 0 & 1 & 2 & 3 & 4 & 4 \\ \hline
 \times -1 & \times -1 & \times -1 & & & \\ \hline
 0 & 0 & 1 & 1 & 2 & 2 \\ \hline
 1 & 1 & 2 & 2 & 3 & 2 \\ \hline
 1 & 2 & 3 & 3 & 2 & 1 \\ \hline
 \end{array} & * & \begin{array}{|c|c|c|} \hline
 -1 & -1 & -1 \\ \hline
 2 & 2 & 2 \\ \hline
 -1 & -1 & -1 \\ \hline
 \end{array} & = & \begin{array}{|c|c|c|c|} \hline
 2 & 3 & 0 & -3 \\ \hline
 -2 & 0 & 4 & 8 \\ \hline
 -5 & -7 & -8 & -8 \\ \hline
 1 & 0 & 2 & 3 \\ \hline
 \end{array} \\
 I & & K & & I * K
 \end{array}$$

Napomena: ovaj konretan primer sadrži bijas, ali je dovoljan za razumevanje način na koji funkcionišu filteri kod konvolucija.

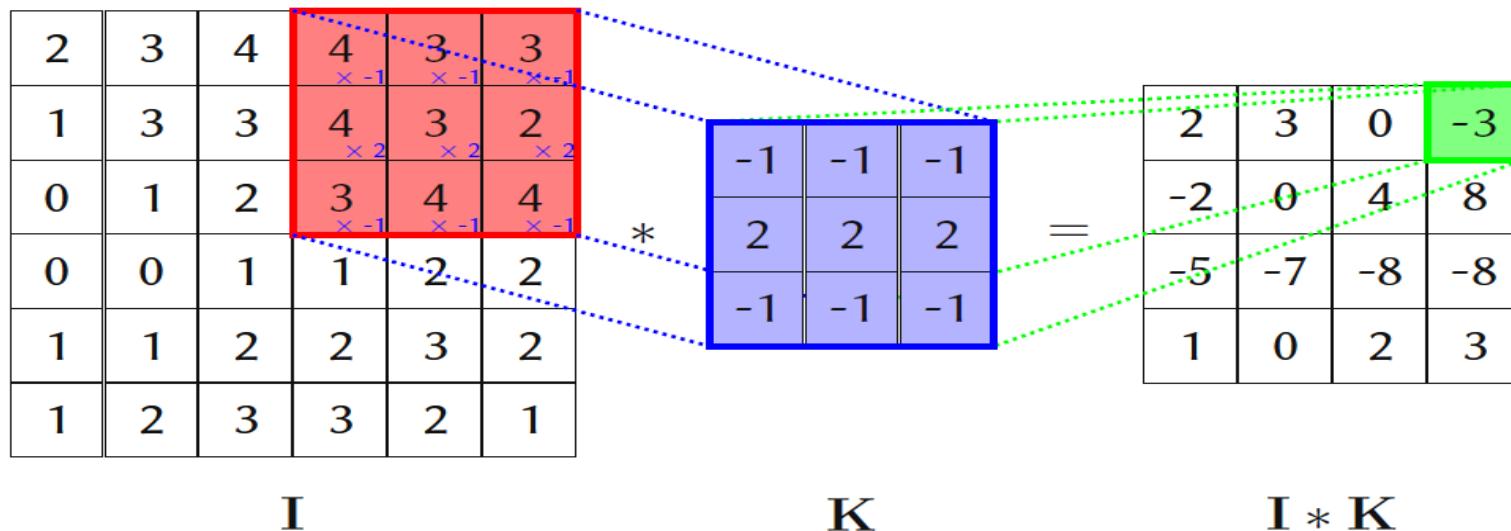
Primer konvolucije



Primer konvolucije



Primer konvolucije



Primer konvolucije

$$\begin{array}{|c|c|c|c|c|c|} \hline
 2 & 3 & 4 & 4 & 3 & 3 \\ \hline
 1 & 3 & 3 & 4 & 3 & 2 \\ \hline
 0 & 1 & 2 & 3 & 4 & 4 \\ \hline
 0 & 0 & 1 & 1 & 2 & 2 \\ \hline
 1 & 1 & 2 & 2 & 3 & 2 \\ \hline
 1 & 2 & 3 & 3 & 2 & 1 \\ \hline
 \end{array}
 \quad *
 \quad
 \begin{array}{|c|c|c|} \hline
 -1 & -1 & -1 \\ \hline
 2 & 2 & 2 \\ \hline
 -1 & -1 & -1 \\ \hline
 \end{array}
 \quad =
 \quad
 \begin{array}{|c|c|c|c|} \hline
 2 & 3 & 0 & -3 \\ \hline
 -2 & 0 & 4 & 8 \\ \hline
 -5 & -7 & -8 & -8 \\ \hline
 1 & 0 & 2 & 3 \\ \hline
 \end{array}$$

I

K

$I * K$

The diagram illustrates a convolution operation. The input matrix I (6x6) has values: row 1: 2, 3, 4, 4, 3, 3; row 2: 1, 3, 3, 4, 3, 2; row 3: 0, 1, 2, 3, 4, 4; row 4: 0, 0, 1, 1, 2, 2; row 5: 1, 1, 2, 2, 3, 2; row 6: 1, 2, 3, 3, 2, 1. A 3x3 kernel K (blue) with values: -1, -1, -1; 2, 2, 2; -1, -1, -1 is applied. The result matrix $I * K$ (4x4) has values: row 1: 2, 3, 0, -3; row 2: -2, 0, 4, 8; row 3: -5, -7, -8, -8; row 4: 1, 0, 2, 3. The result cell at position (2,2) is highlighted in green.

Primer konvolucije

$$\begin{array}{|c|c|c|c|c|c|} \hline
 2 & 3 & 4 & 4 & 3 & 3 \\ \hline
 1 & 3 & 3 & 4 & 3 & 2 \\ \hline
 0 & 1 & 2 & 3 & 4 & 4 \\ \hline
 0 & 0 & 1 & 1 & 2 & 2 \\ \hline
 1 & 1 & 2 & 2 & 3 & 2 \\ \hline
 1 & 2 & 3 & 3 & 2 & 1 \\ \hline
 \end{array}
 \quad *
 \quad
 \begin{array}{|c|c|c|} \hline
 -1 & -1 & -1 \\ \hline
 2 & 2 & 2 \\ \hline
 -1 & -1 & -1 \\ \hline
 \end{array}
 \quad =
 \quad
 \begin{array}{|c|c|c|c|} \hline
 2 & 3 & 0 & -3 \\ \hline
 -2 & 0 & 4 & 8 \\ \hline
 -5 & -7 & -8 & -8 \\ \hline
 1 & 0 & 2 & 3 \\ \hline
 \end{array}$$

I

K

$I * K$

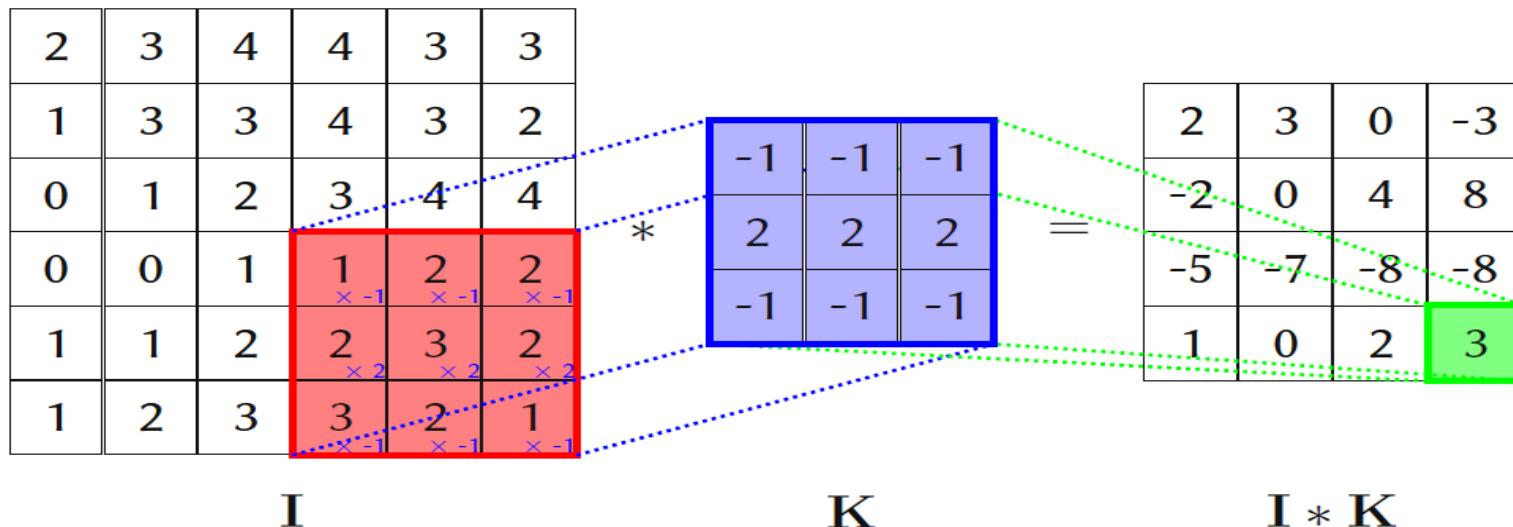
The diagram illustrates a convolution operation. The input matrix I (6x6) has values: row 1: 2, 3, 4, 4, 3, 3; row 2: 1, 3, 3, 4, 3, 2; row 3: 0, 1, 2, 3, 4, 4; row 4: 0, 0, 1, 1, 2, 2; row 5: 1, 1, 2, 2, 3, 2; row 6: 1, 2, 3, 3, 2, 1. The kernel K (3x3) has values: top row: -1, -1, -1; middle row: 2, 2, 2; bottom row: -1, -1, -1. The result matrix $I * K$ (4x4) has values: row 1: 2, 3, 0, -3; row 2: -2, 0, 4, 8; row 3: -5, -7, -8, -8; row 4: 1, 0, 2, 3. A green box highlights the value 0 in the result matrix at position (2,2).

Primer konvolucije

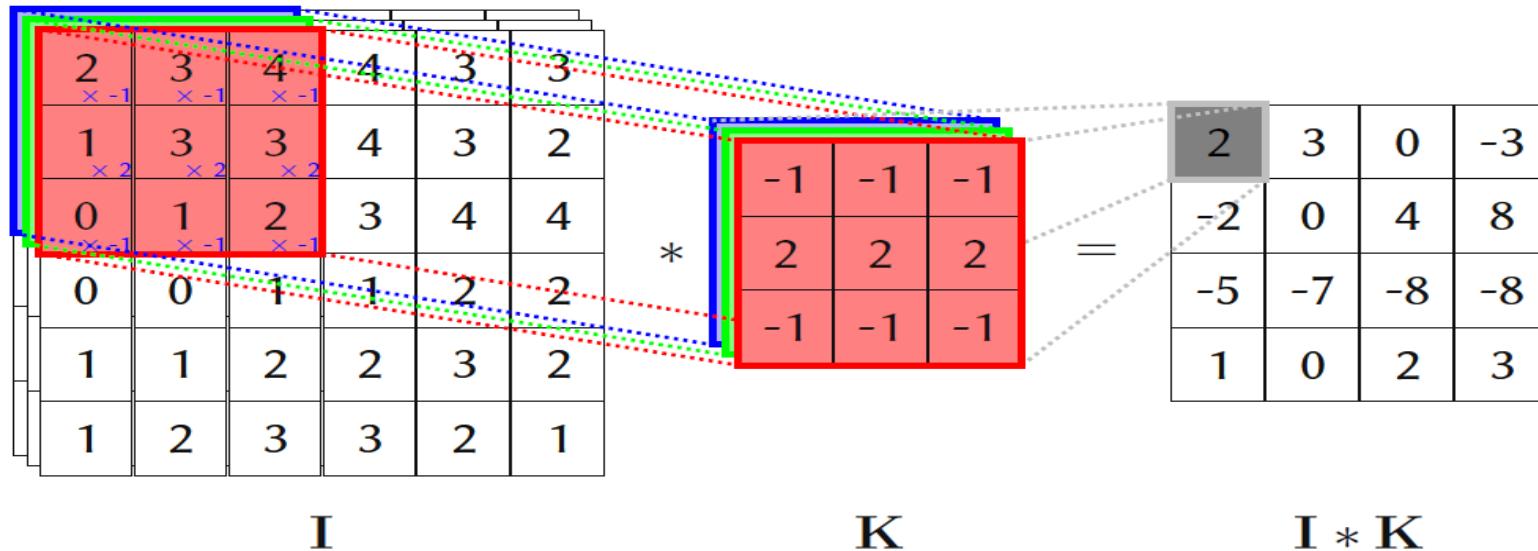
$$\begin{array}{|c|c|c|c|c|c|c|} \hline
 2 & 3 & 4 & 4 & 3 & 3 \\ \hline
 1 & 3 & 3 & 4 & 3 & 2 \\ \hline
 & & \times -1 & \times -1 & \times -1 & \\ \hline
 0 & 1 & 2 & 3 & 4 & 4 \\ \hline
 & & \times 2 & \times 2 & \times 2 & \\ \hline
 0 & 0 & 1 & 1 & 2 & 2 \\ \hline
 & & \times -1 & \times -1 & \times -1 & \\ \hline
 1 & 1 & 2 & 2 & 3 & 2 \\ \hline
 1 & 2 & 3 & 3 & 2 & 1 \\ \hline
 \end{array}
 \quad *
 \quad
 \begin{array}{|c|c|c|} \hline
 -1 & -1 & -1 \\ \hline
 2 & 2 & 2 \\ \hline
 -1 & -1 & -1 \\ \hline
 \end{array}
 \quad =
 \quad
 \begin{array}{|c|c|c|c|} \hline
 2 & 3 & 0 & -3 \\ \hline
 -2 & 0 & 4 & 8 \\ \hline
 -5 & 7 & -8 & -8 \\ \hline
 1 & 0 & 2 & 3 \\ \hline
 \end{array}$$

I **K** **$I * K$**

Primer konvolucije



Konvolucija sa bojom/kanalima



0	0	0	0	0	0	...
0	156	155	156	158	158	...
0	153	154	157	159	159	...
0	149	151	155	158	159	...
0	146	146	149	153	158	...
0	145	143	143	148	158	...
...

Input Channel #1 (Red)

-1	-1	1
0	1	-1
0	1	1

Kernel Channel #1



308

+

0	0	0	0	0	0	...
0	167	166	167	169	169	...
0	164	165	168	170	170	...
0	160	162	166	169	170	...
0	156	156	159	163	168	...
0	155	153	153	158	168	...
...

Input Channel #2 (Green)

1	0	0
1	-1	-1
1	0	-1

Kernel Channel #2



-498

0	0	0	0	0	0	...
0	163	162	163	165	165	...
0	160	161	164	166	166	...
0	156	158	162	165	166	...
0	155	155	158	162	167	...
0	154	152	152	157	167	...
...

Input Channel #3 (Blue)

0	1	1
0	1	0
1	-1	1

Kernel Channel #3



164

+

+ 1 = -25

Bias = 1

-25				...
				...
				...
				...
...

Konvolucija

Šta je ustvari konvolucija? To je filter koji nam služi za ekstrakciju karakteristika slike.



Konvolucija

Na primer na ovoj slici primenom kovolucije dobijamo ivice.



Konvolucija

Zajednica ljudi koji se bave Kompjuterskom Vizijom je ranije morala ručno da razvija konvolucije



Konvolucija

Konvolutivne mreže same uče konvolucije! To je poenta Deep Learning!



Primer konvolucije sa za detekciju izvica

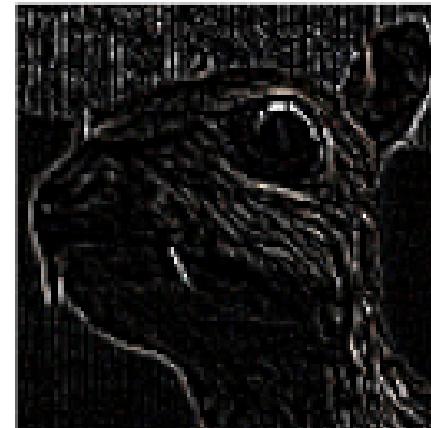
Input image



Convolution
Kernel

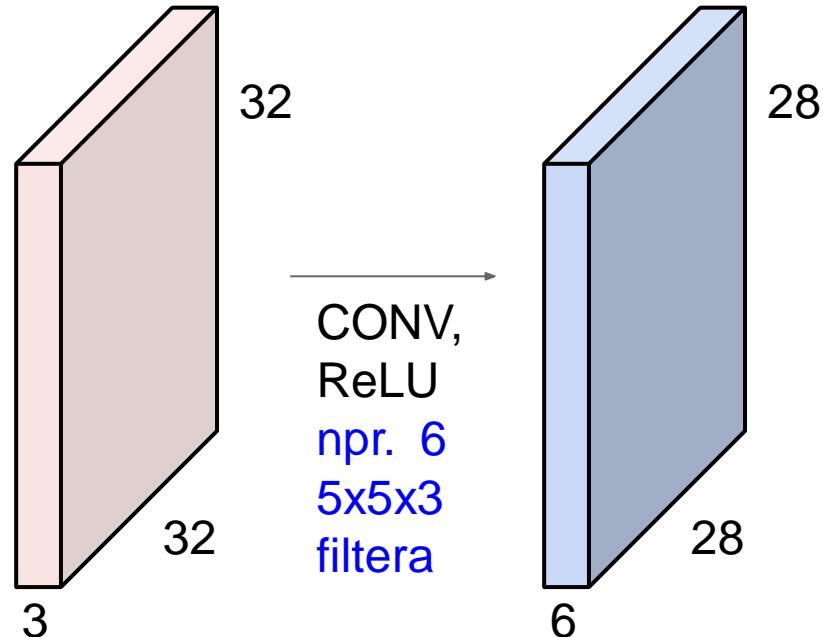
$$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$$

Feature map

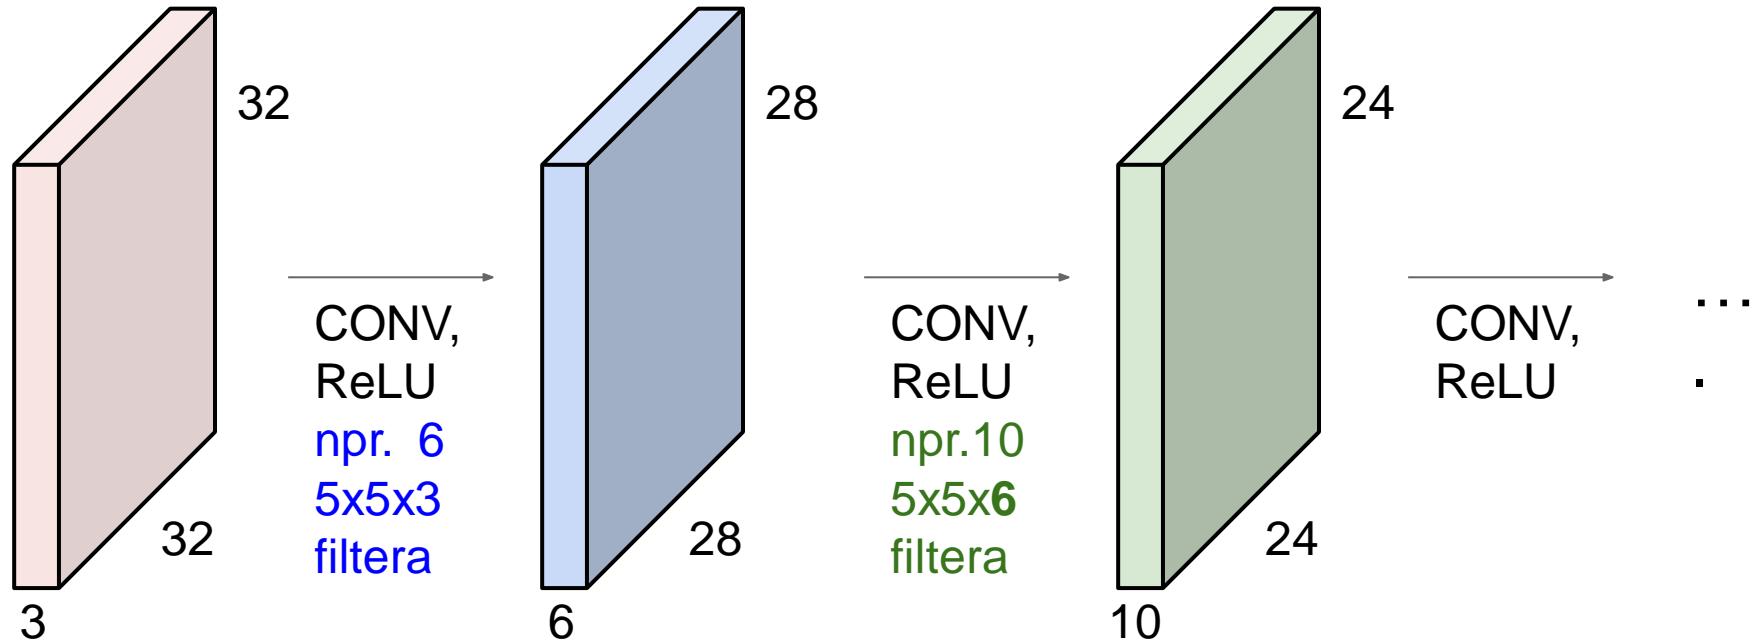


Vrednosti matrice je ručno kreirao ekspert.
Dok ih kovolutivna mreža sama uči!

Konvolutivna mreža je niz Konvolutivnih Slojeva, između kojih su aktivacione funkcije

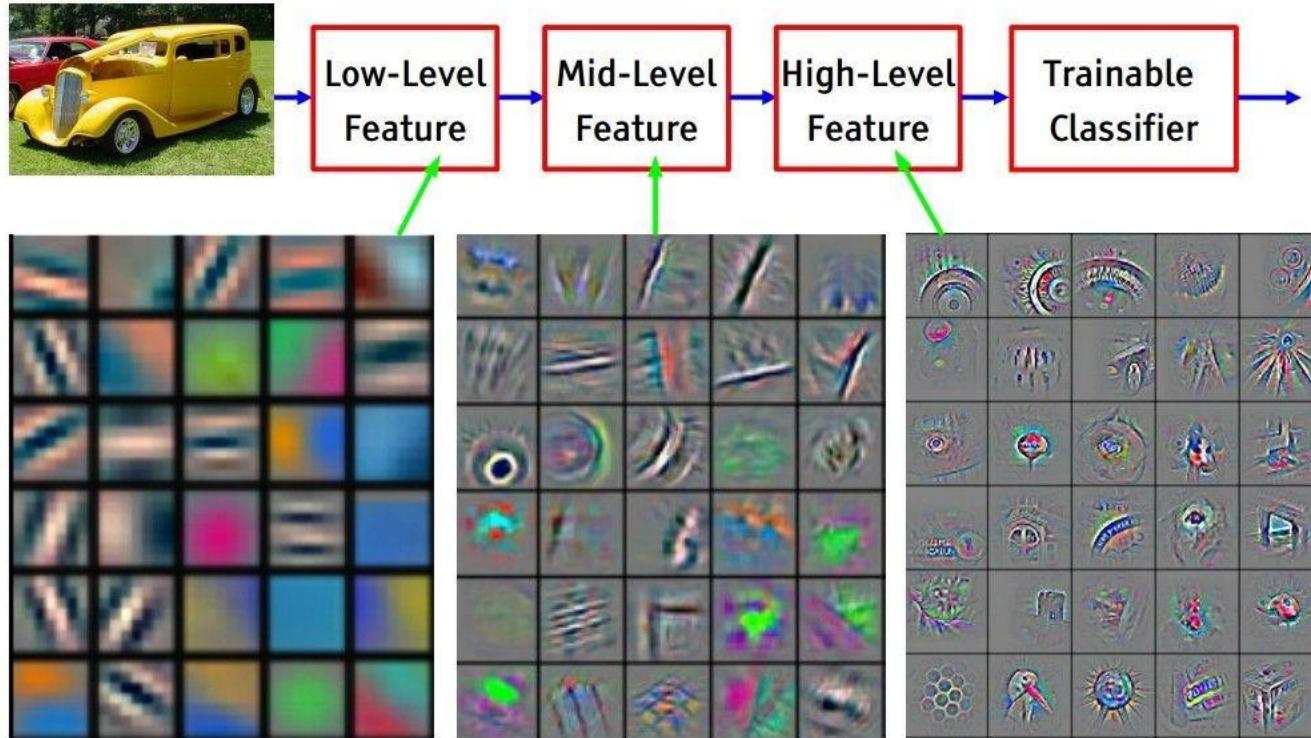


Konvolutivna mreža je niz Konvolutivnih Slojeva, između kojih su aktivacione funkcije



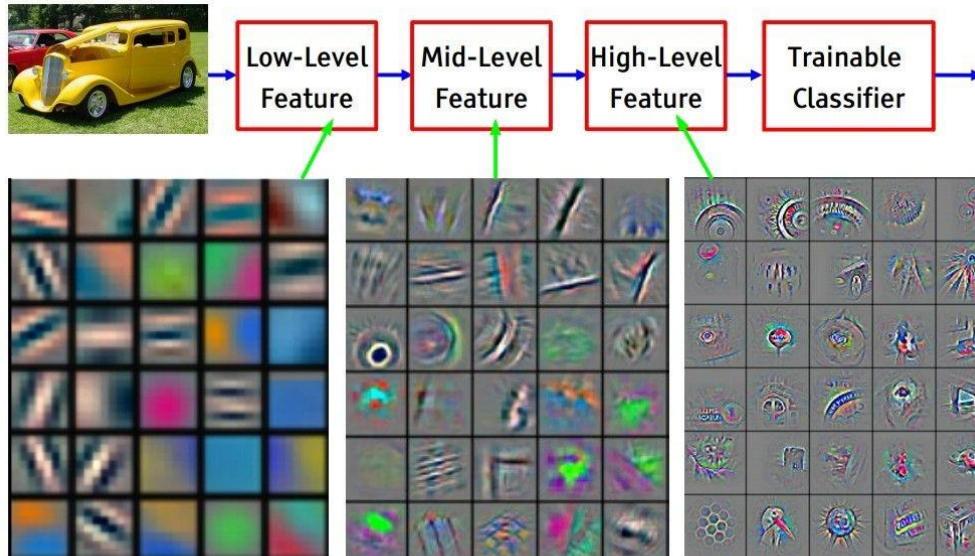
Kako izgledaju slojevi

[From recent Yann LeCun slides]



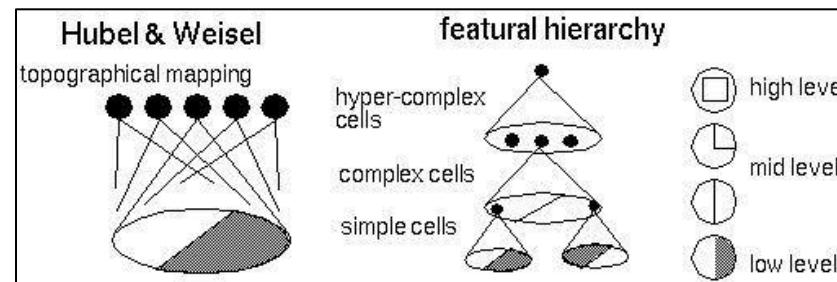
Feature visualization of convolutional net trained on ImageNet from [Zeiler & Fergus 2013]

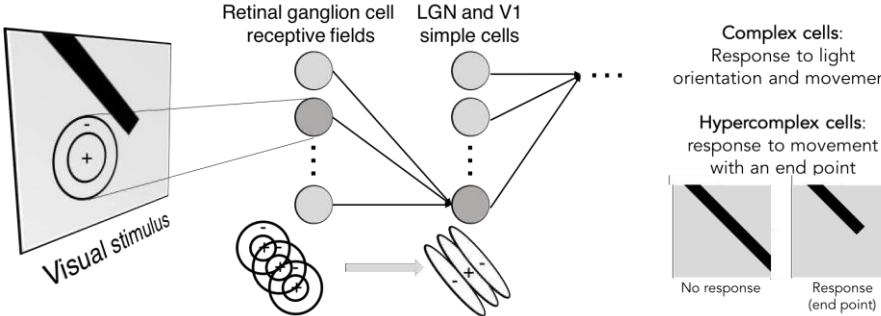
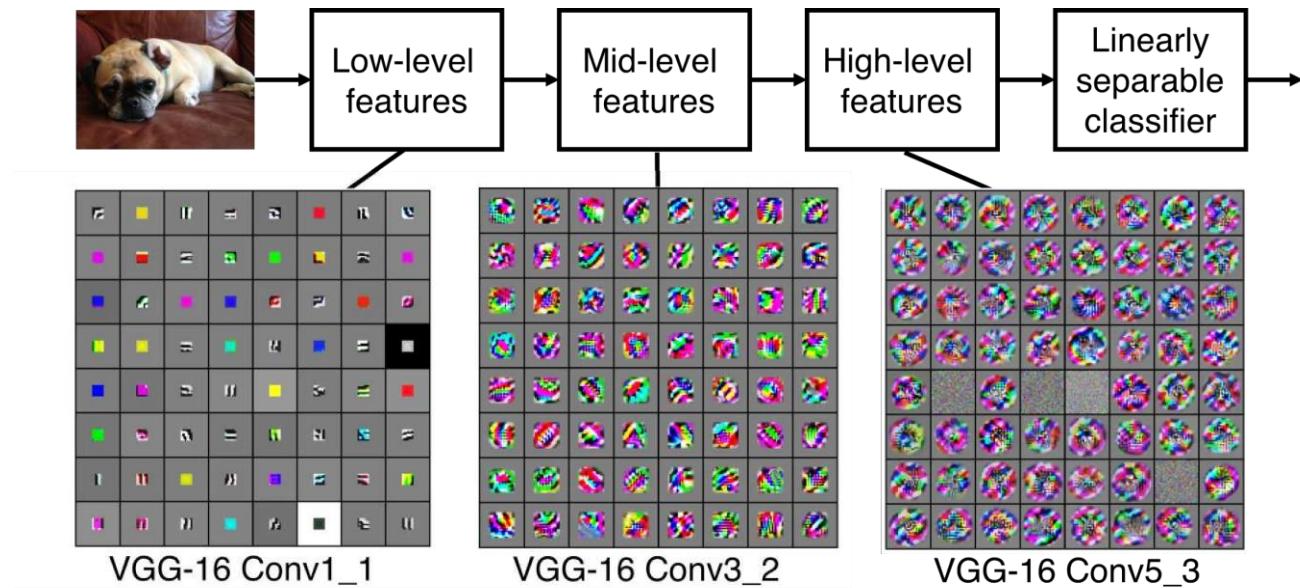
Kako izgledaju slojevi



[From recent Yann LeCun slides]

Feature visualization of convolutional net trained on ImageNet from [Zeiler & Fergus 2013]

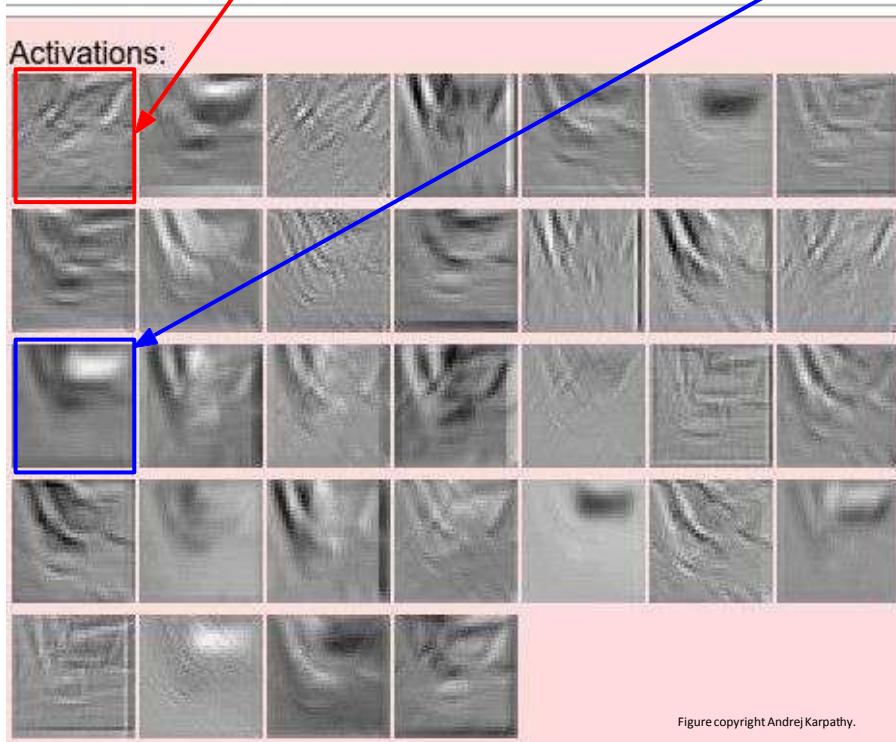




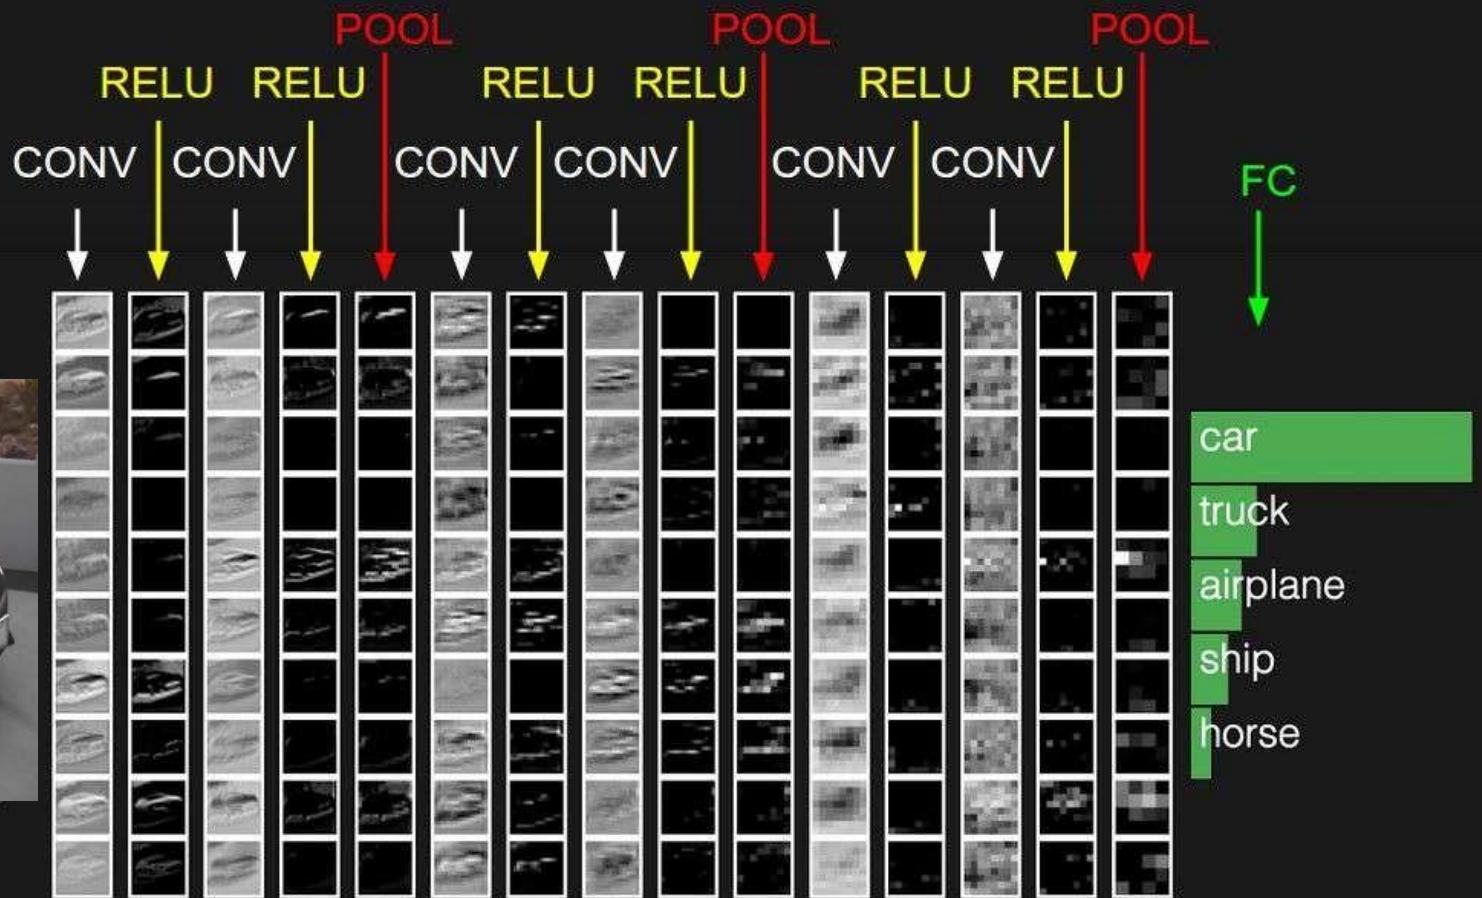


jedan filter =>
jedna aktivaciona mapa

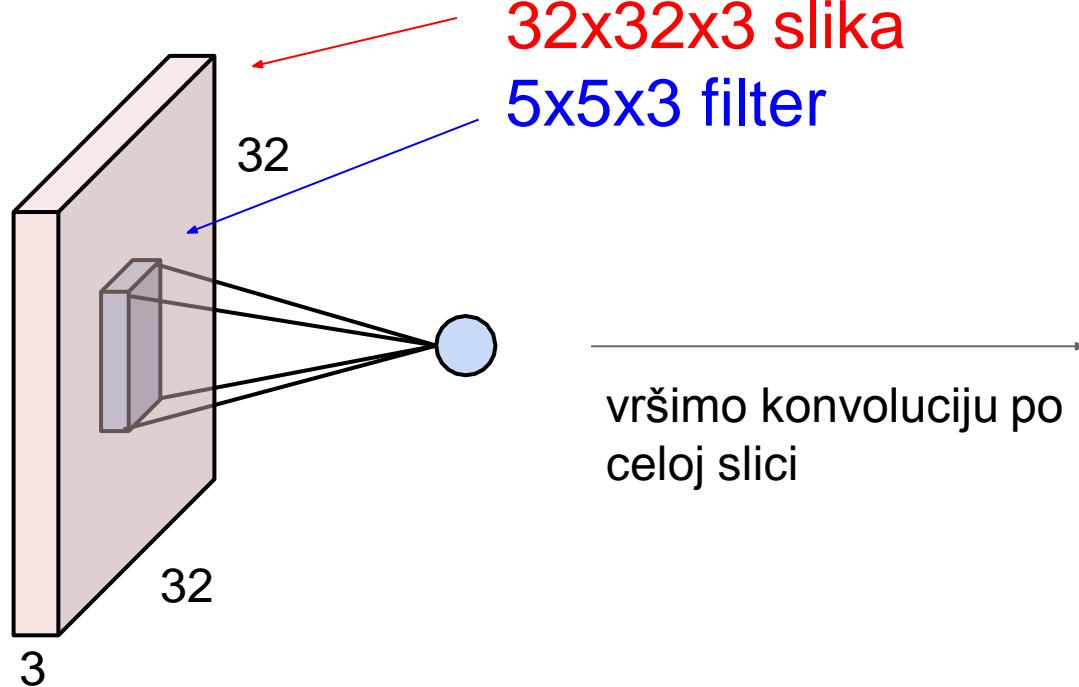
primer 5x5 filteri
(32 ukupno)



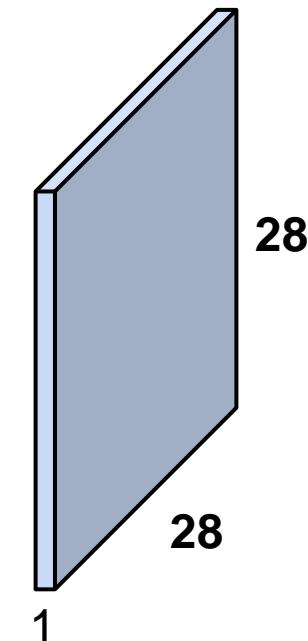
Kako tipično izgleda Konvolutivna Mreža



Detaljnije o dimenzijama:

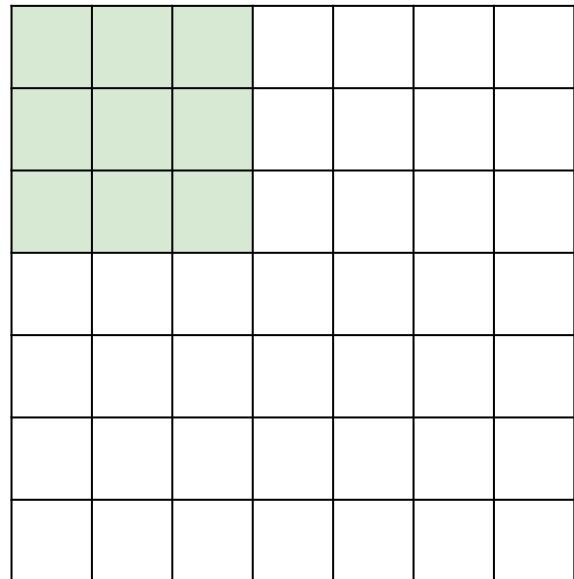


aktivaciona mapa



Detaljnije o dimenzijama:

7

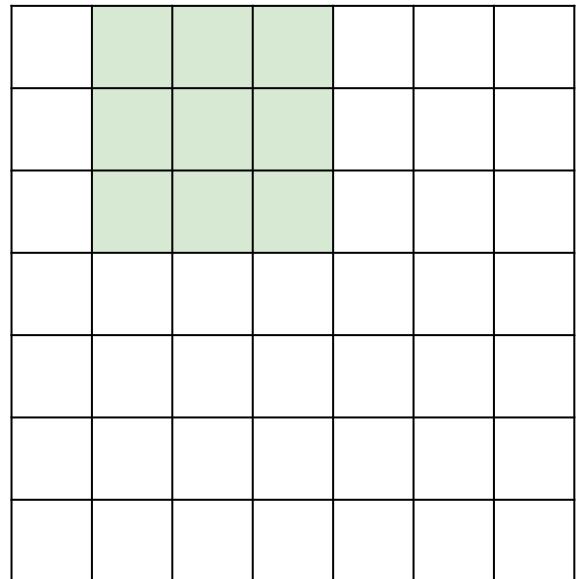


7x7 ulaz
3x3 filter

7

Detaljnije o dimenzijama:

7

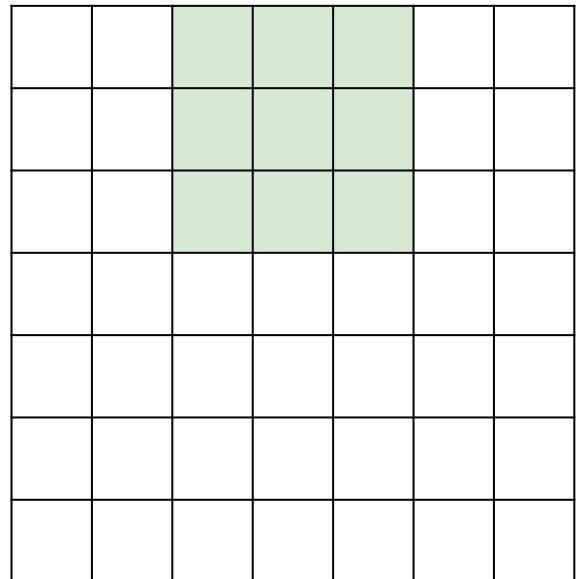


7x7 ulaz
3x3 filter

7

Detaljnije o dimenzijama:

7

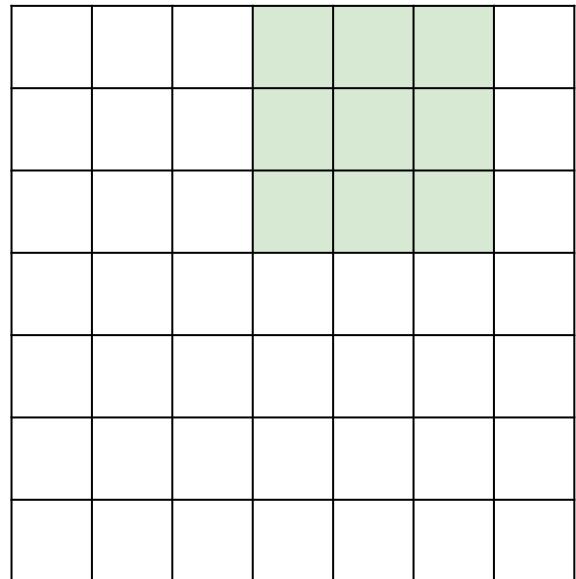


7x7 ulaz
3x3 filter

7

Detaljnije o dimenzijama:

7

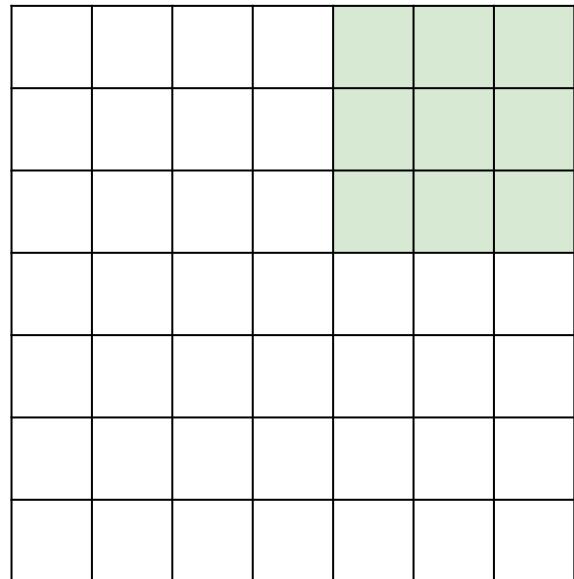


7x7 ulaz
3x3 filter

7

Detaljnije o dimenzijama:

7



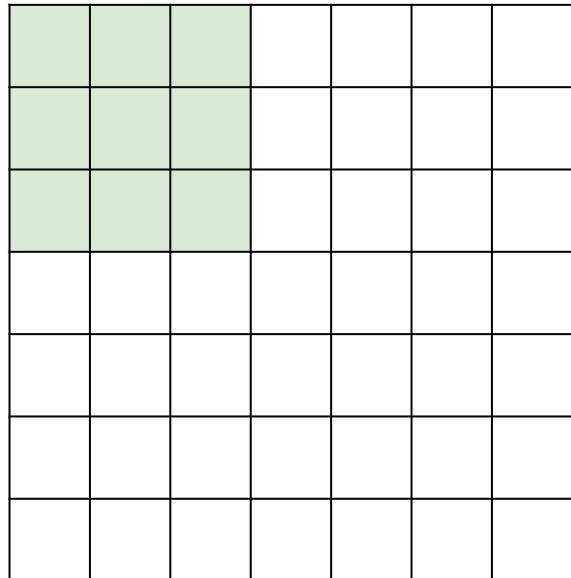
7x7 ulaz
3x3 filter

=> 5x5 izlaz

7

Detaljnije o dimenzijama:

7



7x7 ulza

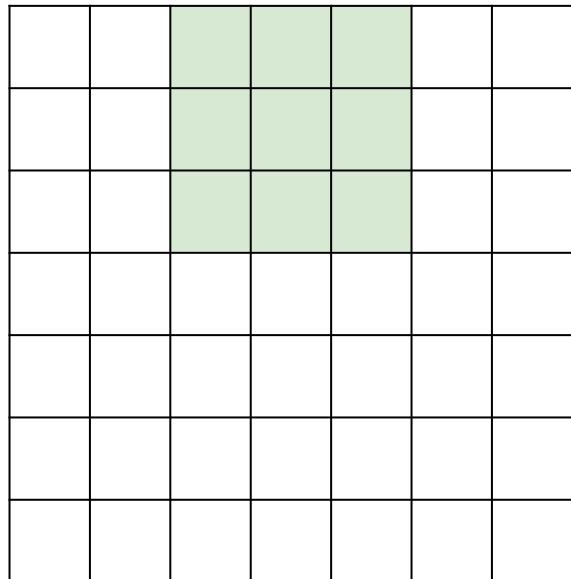
3x3 filter

primenjen sa **korakom (stride) 2**

7

Detaljnije o dimenzijama:

7



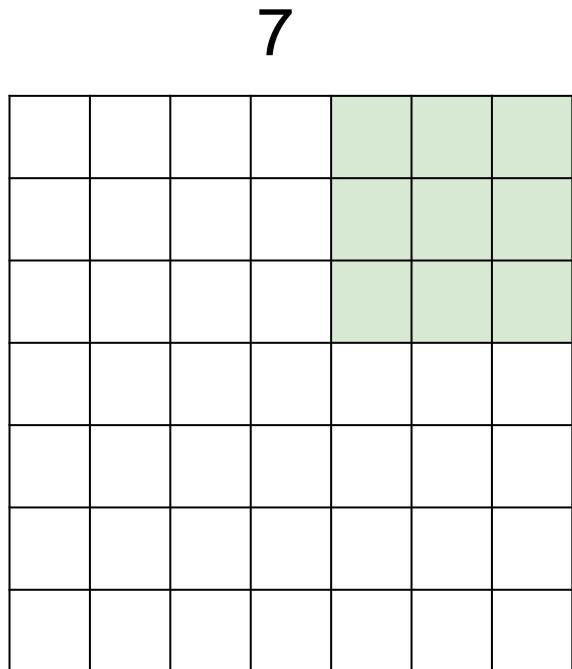
7x7 ulza

3x3 filter

primenjen sa **korakom (stride) 2**

7

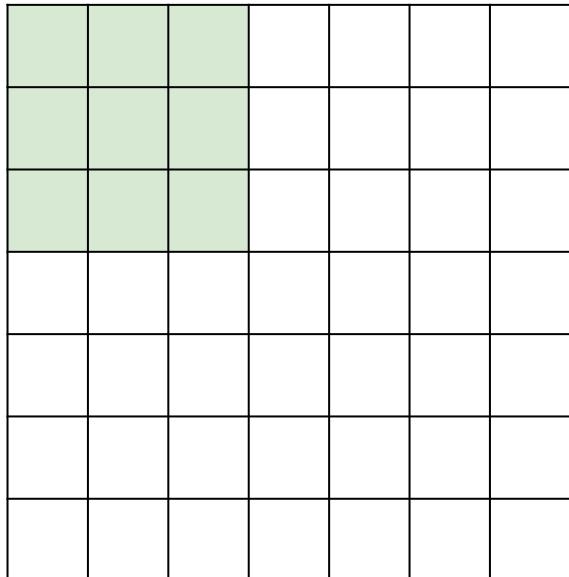
Detaljnije o dimenzijama:



7x7 ulza
3x3 filter
primenjen sa **korakom (stride) 2**
=> 3x3 izlaz!

Detaljnije o dimenzijama:

7



7x7 ulaz

3x3 filter

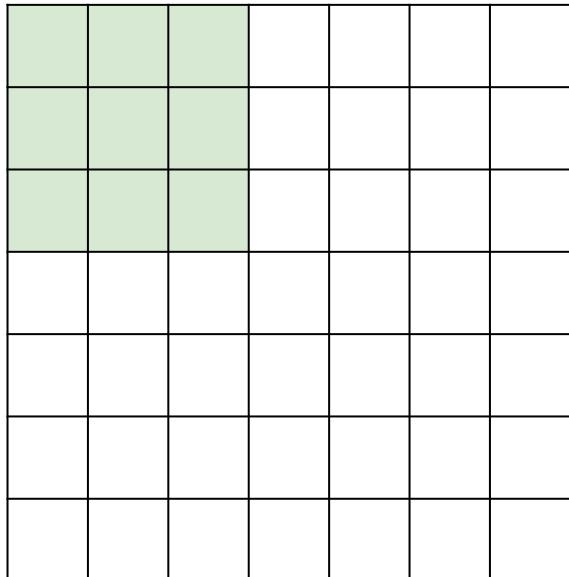
primenjen sa **korakom (stride) 3**

Da li je primena moguća?

7

Detaljnije o dimenzijama:

7



7x7 ulaz

3x3 filter

primenjen sa **korakom (stride) 3**

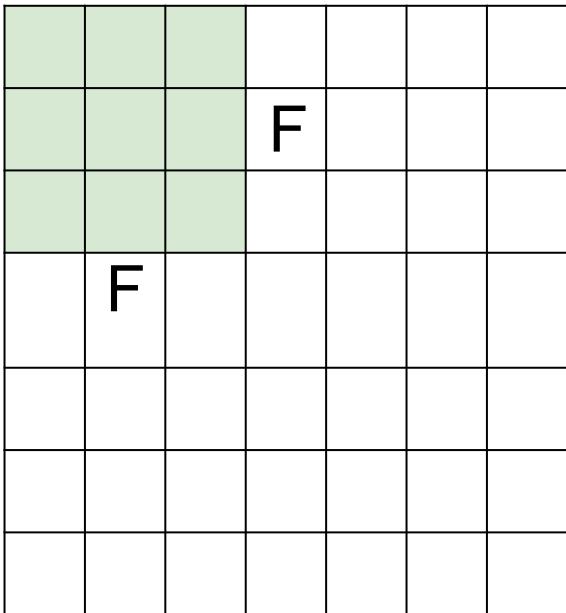
Da li je primena moguća?

7

Ne, filter se ne uklapa!

ne možemo da применимо 3x3 filter
na 7x7 ulaz sa korakom 3.

N



Dimenzionalnost izlaza:
((N - F) / korak) + 1

N

npr. N = 7, F = 3:

$$\text{korak 1} \Rightarrow (7 - 3)/1 + 1 = 5$$

$$\text{korak 2} \Rightarrow (7 - 3)/2 + 1 = 3$$

$$\text{korak 3} \Rightarrow (7 - 3)/3 + 1 = 2.33$$

– ne dobijamo ceo broj

U praksi: Dodajemo nule u okvir slike

0	0	0	0	0	0		
0							
0							
0							
0							

npr. ulaz 7×7

3×3 filter, primenjen sa korakom 1

dodajemo okvir od 1 piksela => koje je su dimenzije izlaza?

$$\begin{aligned} &(\text{podsetite se:}) \\ &\textcolor{blue}{(N - F) / \text{korak} + 1} \end{aligned}$$

U praksi: Dodajemo nule u okvir slike

0	0	0	0	0	0		
0							
0							
0							
0							

npr. ulaz 7x7

3x3 filter, primenjen sa **korakom 1**

dodajemo okvir od 1 piksela => koje je su dimenzije izlaza?

Izlaz je 7X7! $(7+2\cdot3)/1 + 1 = 7$

podsetite se:

$(N - F) / \text{korak} + 1$

U praksi: Dodajemo nule u okvir slike

0	0	0	0	0	0		
0							
0							
0							
0							

npr. ulaz 7x7

3x3 filter, primjenjen sa korakom 1

dodajemo okvir od 1 piksela => koje je su dimenzijske izlaza?

Izlaz je 7X7! $(7+2-3)/1 + 1 = 7$

tipično se koriste CONV slojevi sa korakom 1, filteri veličine FxF, dodavanje 0 u okvir dimenzijske $(F-1)/2$.

npr.

F = 3 => dodajemo okvir dim. 1

F = 5 => dodajemo okvir dim. 2

F = 7 => dodajemo okvir dim. 3

Primer konvolucije (sa *padding*-om)

0	0	0	0	0	0	0	0
0	2	3	4	4	3	3	0
0	1	3	3	4	3	2	0
0	0	1	2	3	4	4	0
0	0	0	1	1	2	2	0
0	1	1	2	2	3	2	0
0	1	2	3	3	2	1	0
0	0	0	0	0	0	0	0

\mathbf{I} (*padded*)

$$* \quad \begin{matrix} -1 & -1 & -1 \\ 2 & 2 & 2 \\ -1 & -1 & -1 \end{matrix}$$

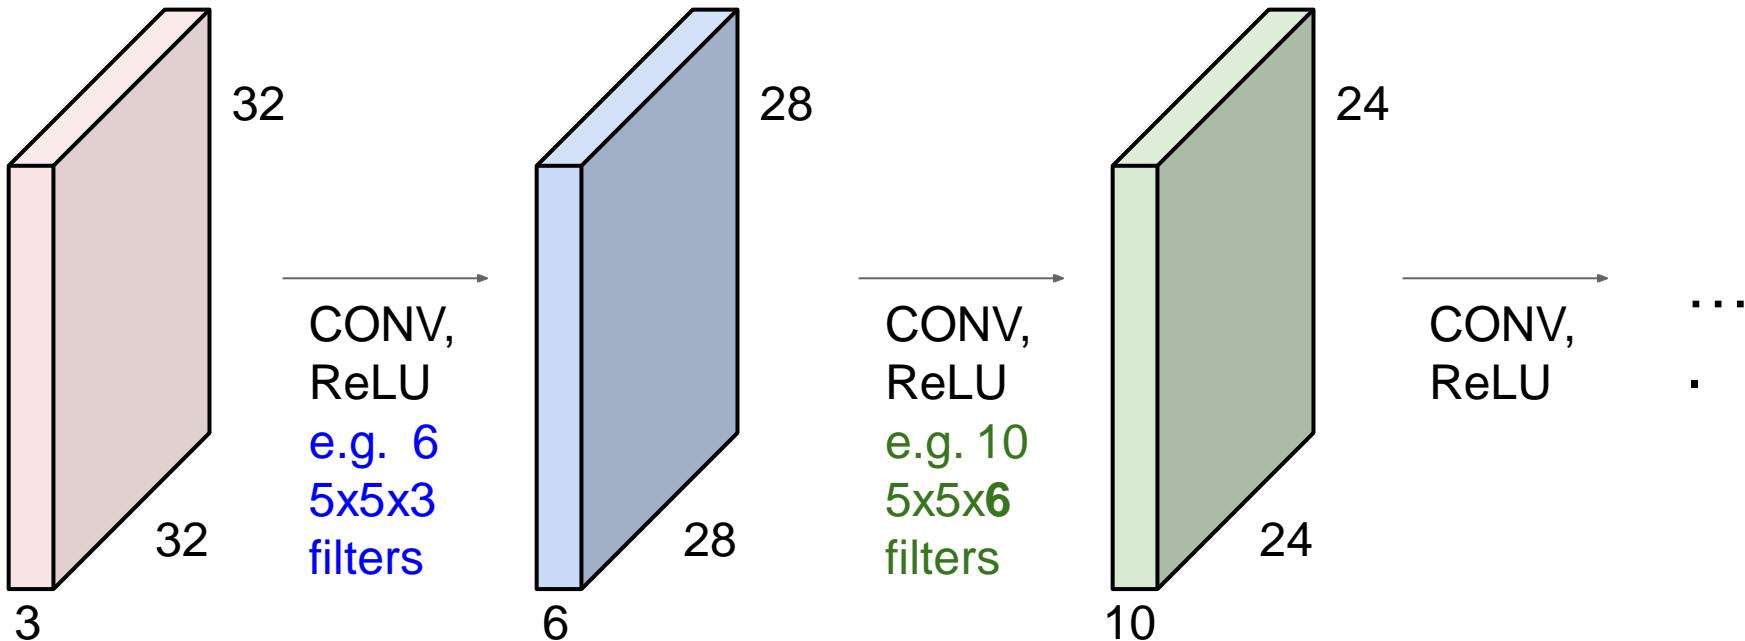
\mathbf{K}

6	11	12	12	11	7
2	2	3	0	-3	-4
-2	-2	0	4	8	7
-3	-5	-7	-8	-8	-5
1	1	0	2	3	3
4	8	11	9	5	1

$\mathbf{I} * \mathbf{K}$

Prijetimo se...

Npr. ulaz 32×32 na kome se u više slojeva primjenjuje filter 5×5 se smanjuje po veličini! ($32 \rightarrow 28 \rightarrow 24 \dots$). Prebrzo smanjivanje nije dobro, gubimo na informacijama, dobijamo loše rezultate. Zato koristimo padding.

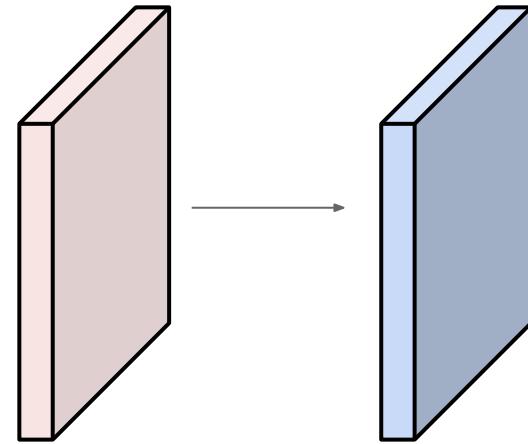


Primer:

Ulag: **32x32x3**

10 5x5x3 filtera sa korakom 1,
i dodavanjem okvira 0 dim. 2

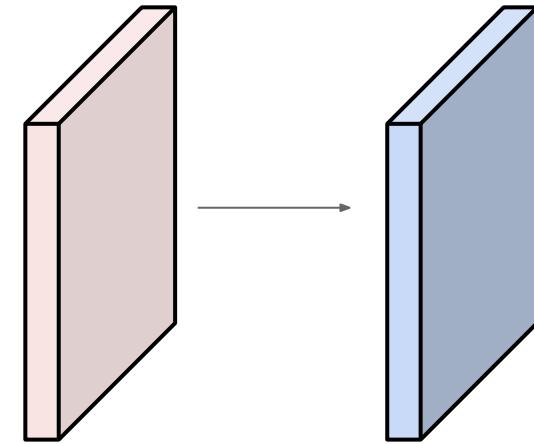
Koja je veličina izlaza: ?



Primer:

Ulag: **32x32x3**

10 **5x5x3** filtera sa korakom **1**,
i dodavanjem okvira 0 dim. **2**



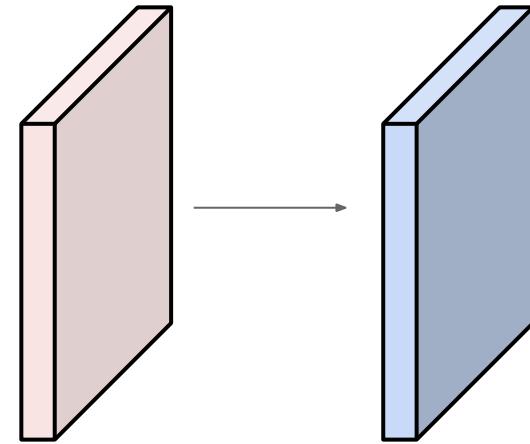
Veličina izlaza:

$(32+2*2-5)/1+1 = 32$ što daje,
32x32x10

Primer:

Ulag: **32x32x3**

10 5x5x3 filter sa korakom 1,
i dodavanjem okvira 0 dim. 2

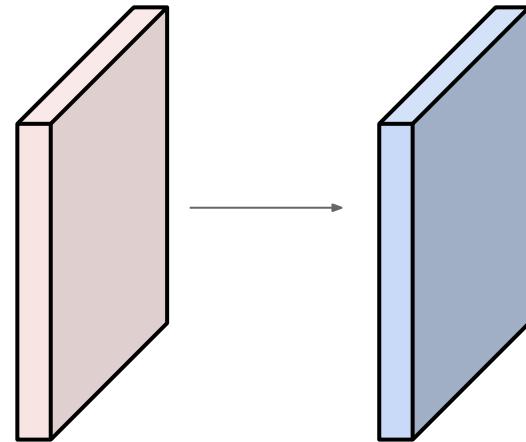


Koji je broj parametara (težina i bijasa)
za obučavanje u ovom sloju?

Primer:

Ulag: **32x32x3**

10 5x5x3 filter sa korakom 1,
i dodavanjem okvira 0 dim. 2



Koji je broj parametara (težina i bijasa) za obučavanje u ovom sloju?

svaki filter ima $5^*5^*3 + 1 = 76$ parametara (+1 za bijas)
 $\Rightarrow 76^*10 = 760$

Summary To summarize, the Conv Layer:

- Accepts a volume of size $W_1 \times H_1 \times D_1$
- Requires four hyperparameters:
 - Number of filters K ,
 - their spatial extent F ,
 - the stride S ,
 - the amount of zero padding P .
- Produces a volume of size $W_2 \times H_2 \times D_2$ where:
 - $W_2 = (W_1 - F + 2P)/S + 1$
 - $H_2 = (H_1 - F + 2P)/S + 1$ (i.e. width and height are computed equally by symmetry)
 - $D_2 = K$
- With parameter sharing, it introduces $F \cdot F \cdot D_1$ weights per filter, for a total of $(F \cdot F \cdot D_1) \cdot K$ weights and K biases.
- In the output volume, the d -th depth slice (of size $W_2 \times H_2$) is the result of performing a valid convolution of the d -th filter over the input volume with a stride of S , and then offset by d -th bias.

Tipična podešavanja:

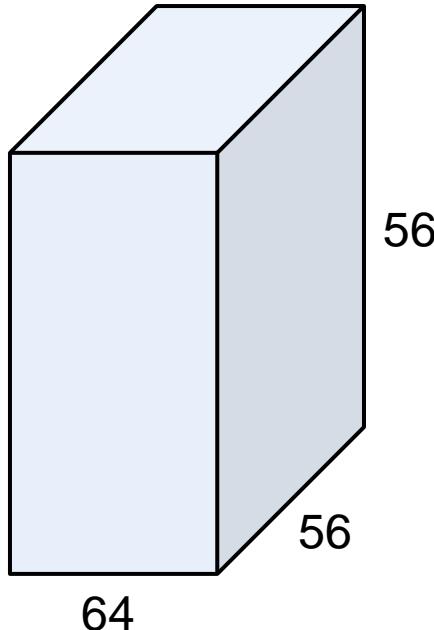
Summary To summarize, the Conv Layer:

- Accepts a volume of size $W_1 \times H_1 \times D_1$
- Requires four hyperparameters:
 - Number of filters K ,
 - their spatial extent F ,
 - the stride S ,
 - the amount of zero padding P .
- Produces a volume of size $W_2 \times H_2 \times D_2$ where:
 - $W_2 = (W_1 - F + 2P)/S + 1$
 - $H_2 = (H_1 - F + 2P)/S + 1$ (i.e. width and height are computed equally by symmetry)
 - $D_2 = K$
- With parameter sharing, it introduces $F \cdot F \cdot D_1$ weights per filter, for a total of $(F \cdot F \cdot D_1) \cdot K$ weights and K biases.
- In the output volume, the d -th depth slice (of size $W_2 \times H_2$) is the result of performing a valid convolution of the d -th filter over the input volume with a stride of S , and then offset by d -th bias.

$K = (\text{steperi 2, npr. } 32, 64, 128, 512)$

- $F = 3, S = 1, P = 1$
- $F = 5, S = 1, P = 2$
- $F = 5, S = 2, P = ?$ (šta god odgovara)
- $F = 1, S = 1, P = 0$

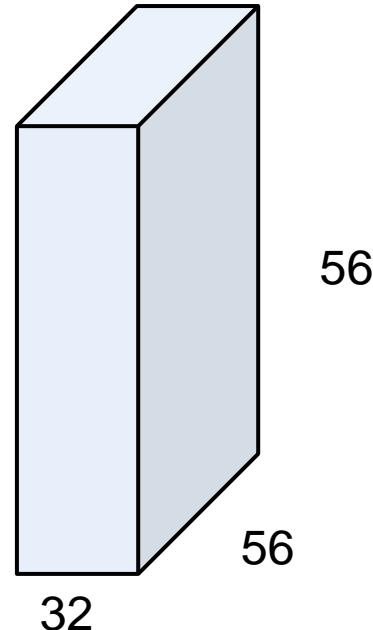
(napomena, 1x1 konvolutivni slojevi imaju smisla)



1x1 CONV
sa 32 filtera

→

(svaki filter je veličine
1x1x64, i računa 64-
dimezionalni skalarni
proizvod)



Primer: CONV sloj u okruženju Torch

SpatialConvolution

```
module = nn.SpatialConvolution(nInputPlane, nOutputPlane, kW, KH, [dW], [dH], [padW], [padH])
```

Applies a 2D convolution over an input image composed of several input planes. The `input` tensor in `forward(input)` is expected to be a 3D tensor (`nInputPlane x height x width`).

The parameters are the following:

- `nInputPlane` : The number of expected input planes in the image given into `forward()`.
- `nOutputPlane` : The number of output planes the convolution layer will produce.
- `kW` : The kernel width of the convolution
- `KH` : The kernel height of the convolution
- `dW` : The step of the convolution in the width dimension. Default is `1`.
- `dH` : The step of the convolution in the height dimension. Default is `1`.
- `padW` : The additional zeros added per width to the input planes. Default is `0`, a good number is `(kW-1)/2`.
- `padH` : The additional zeros added per height to the input planes. Default is `padW`, a good number is `(KH-1)/2`.

Note that depending of the size of your kernel, several (of the last) columns or rows of the input image might be lost. It is up to the user to add proper padding in images.

If the input image is a 3D tensor `nInputPlane x height x width`, the output image size will be `nOutputPlane x oheight x owidth` where

```
owidth = floor((width + 2*padW - kW) / dW + 1)
oheight = floor((height + 2*padH - KH) / dH + 1)
```

[Torch](#) is licensed under [BSD 3-clause](#).

Summary. To summarize, the Conv Layer:

- Accepts a volume of size $W_1 \times H_1 \times D_1$
- Requires four hyperparameters:
 - Number of filters K ,
 - their spatial extent F ,
 - the stride S ,
 - the amount of zero padding P .

Primer: CONV sloj u okruženju *Caffe*

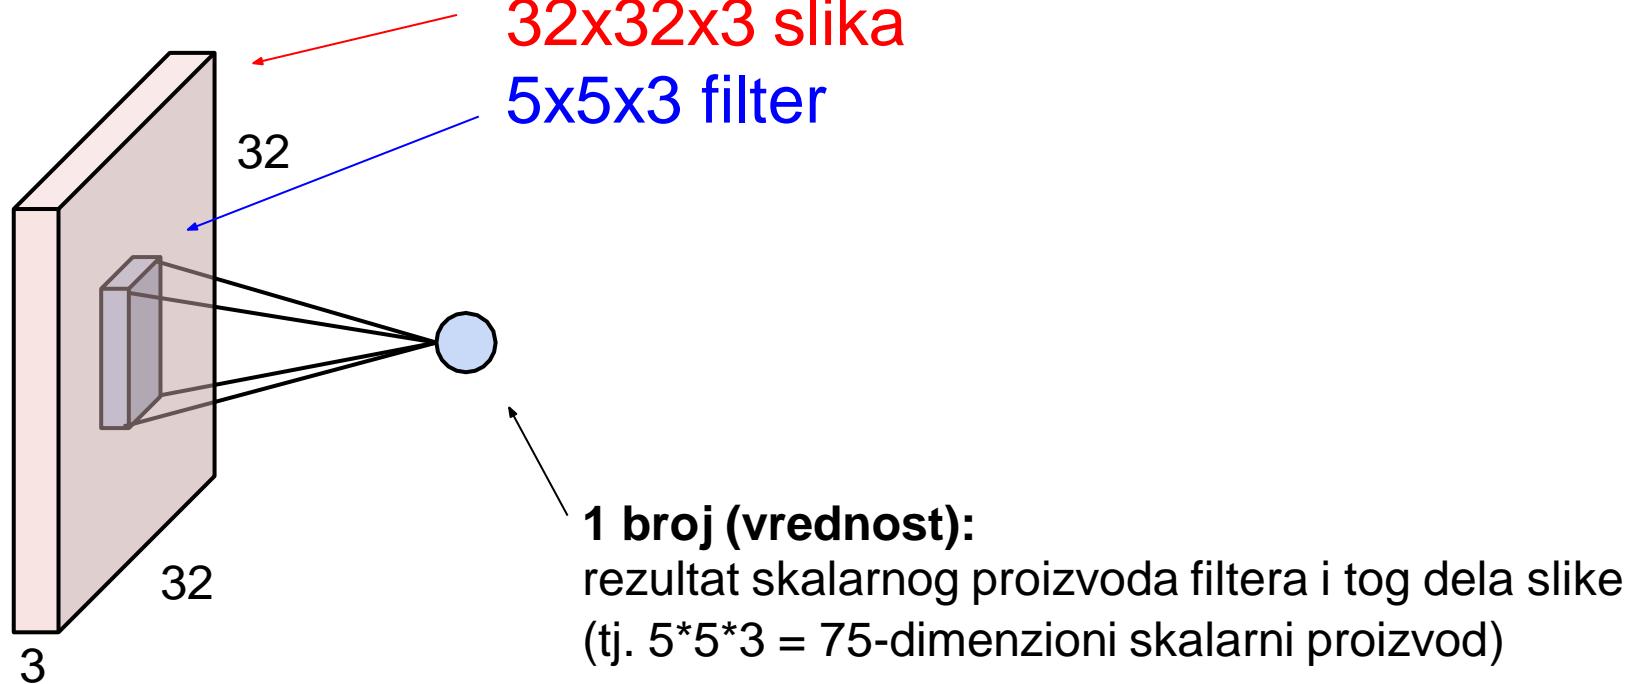
Summary. To summarize, the Conv Layer:

- Accepts a volume of size $W_1 \times H_1 \times D_1$
- Requires four hyperparameters:
 - Number of filters K ,
 - their spatial extent F ,
 - the stride S ,
 - the amount of zero padding P .

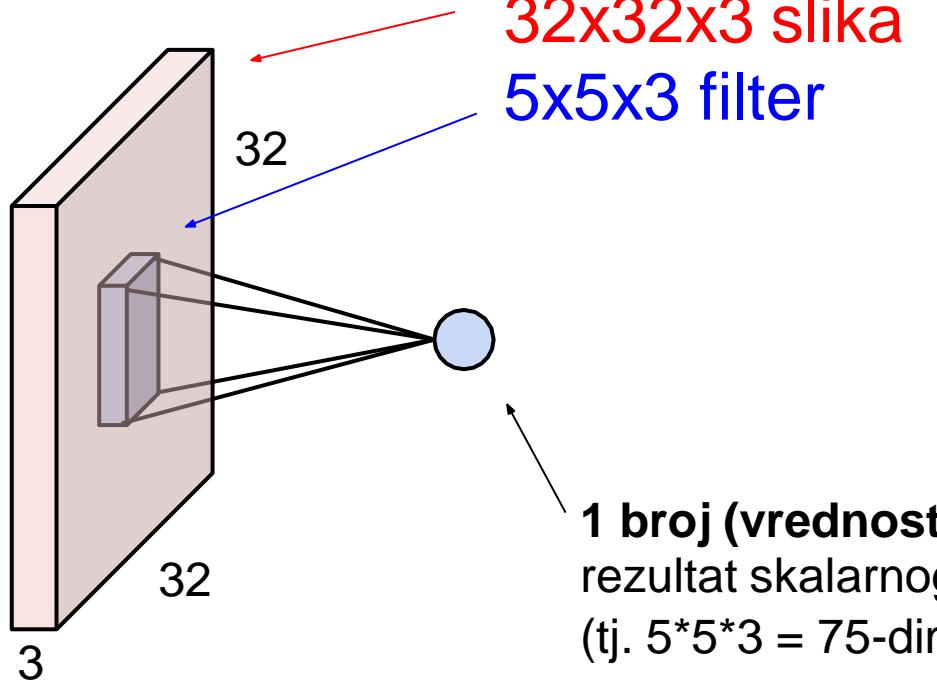
```
layer {
    name: "conv1"
    type: "Convolution"
    bottom: "data"
    top: "conv1"
    # learning rate and decay multipliers for the filters
    param { lr_mult: 1 decay_mult: 1 }
    # learning rate and decay multipliers for the biases
    param { lr_mult: 2 decay_mult: 0 }
    convolution_param {
        num_output: 96      # learn 96 filters
        kernel_size: 11     # each filter is 11x11
        stride: 4           # step 4 pixels between each filter application
        weight_filler {
            type: "gaussian" # initialize the filters from a Gaussian
            std: 0.01          # distribution with stdev 0.01 (default mean: 0)
        }
        bias_filler {
            type: "constant" # initialize the biases to zero (0)
            value: 0
        }
    }
}
```

[Caffe](#) is licensed under [BSD 2-Clause](#).

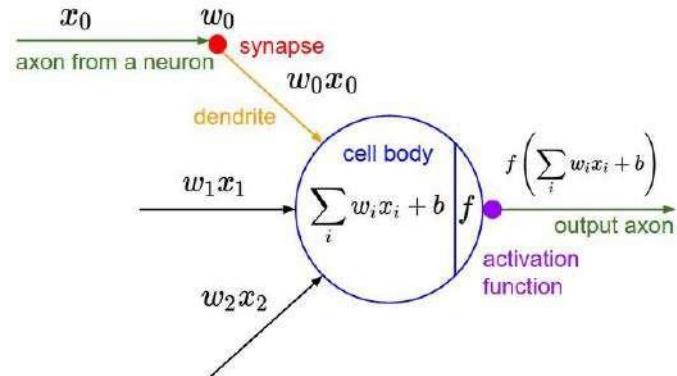
Mozak/neuron pogled na CONV Sloj



Mozak/neuron pogled na CONV Sloj

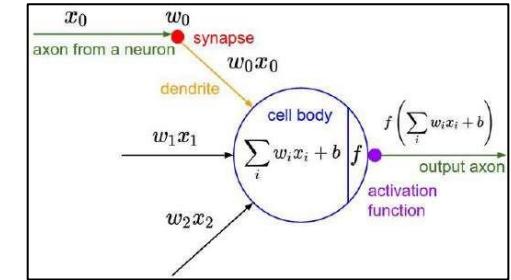
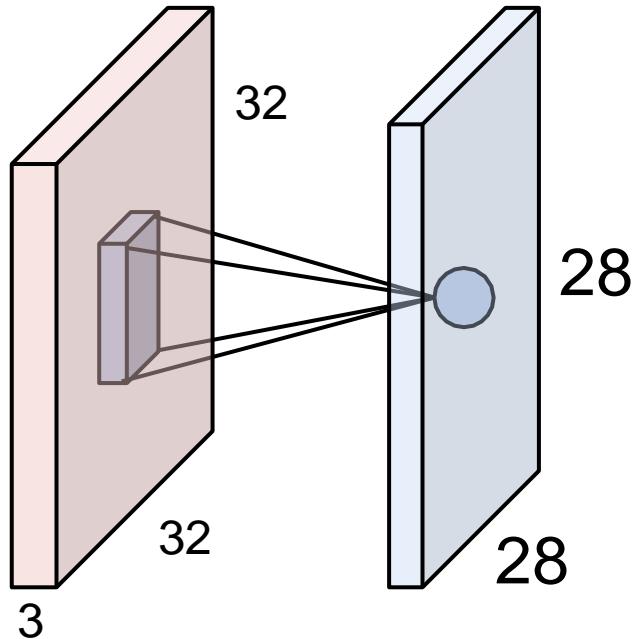


1 broj (vrednost):
rezultat skalarnog proizvoda filtera i tog dela slike
(tj. $5 \times 5 \times 3 = 75$ -dimenzioni skalarni proizvod)



to je samo neuron koji je povezan sa samo jednim delom slike...

Mozak/neuron pogled na CONV Sloj

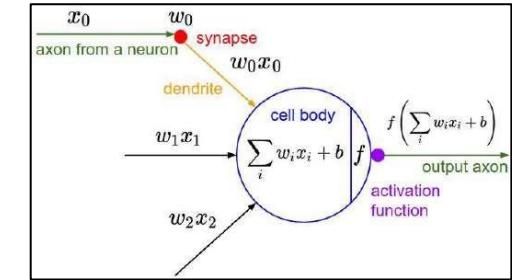
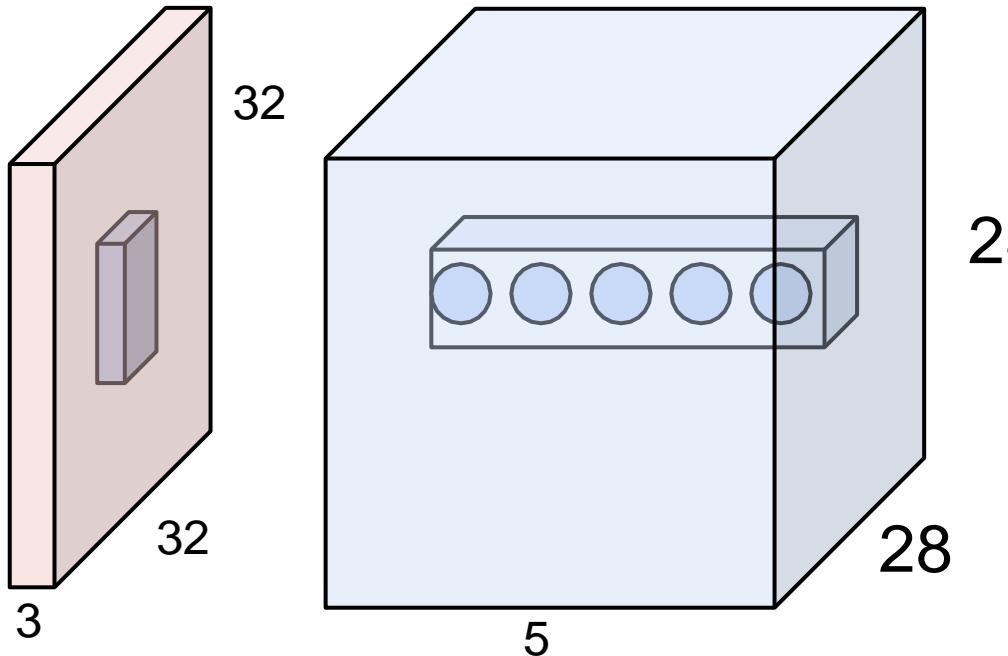


Aktivaciona mapa dimezionalnosti 28×28 je sloj izlaza neurona:

1. Svaki je povezan sa malim delom ulaza
2. Svi dele iste parametre

“5x5 filter” -> “vidno polje dimenzionalnosti 5x5”

Mozak/neuron pogled na CONV Sloj



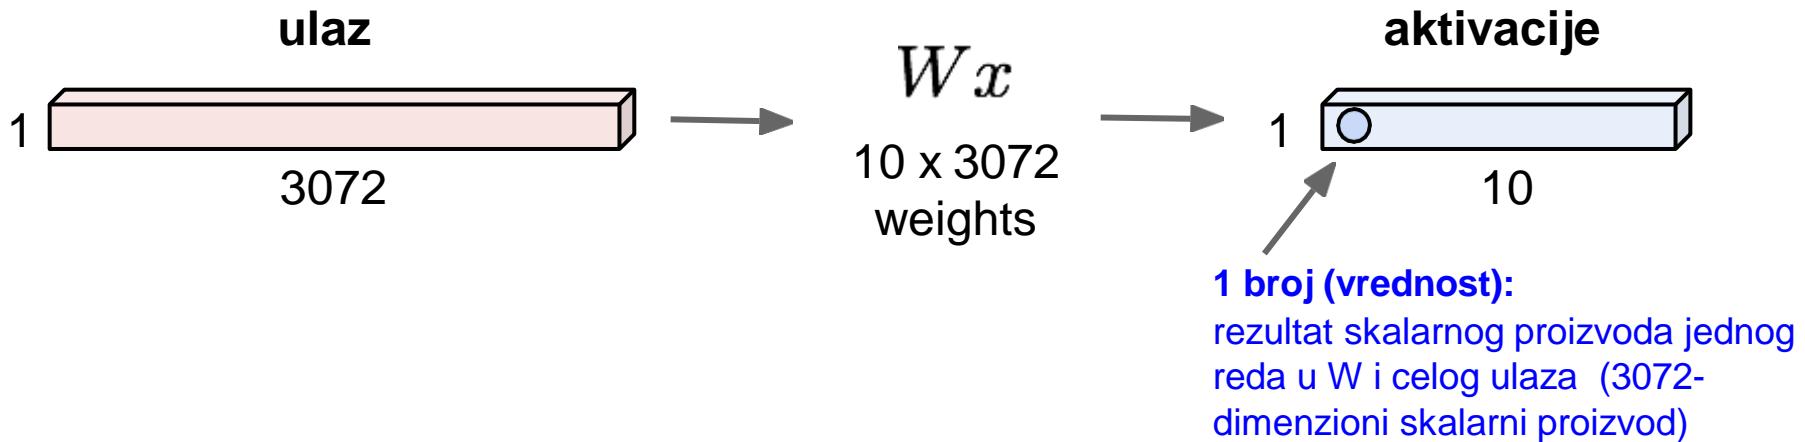
Npr. 5 filtera,
CONV sloj se sastoji od
neurona složenih u mrežu
($28 \times 28 \times 5$)

Imamo 5 različitih neurona koji
su povezani sa istim delom
ulaza

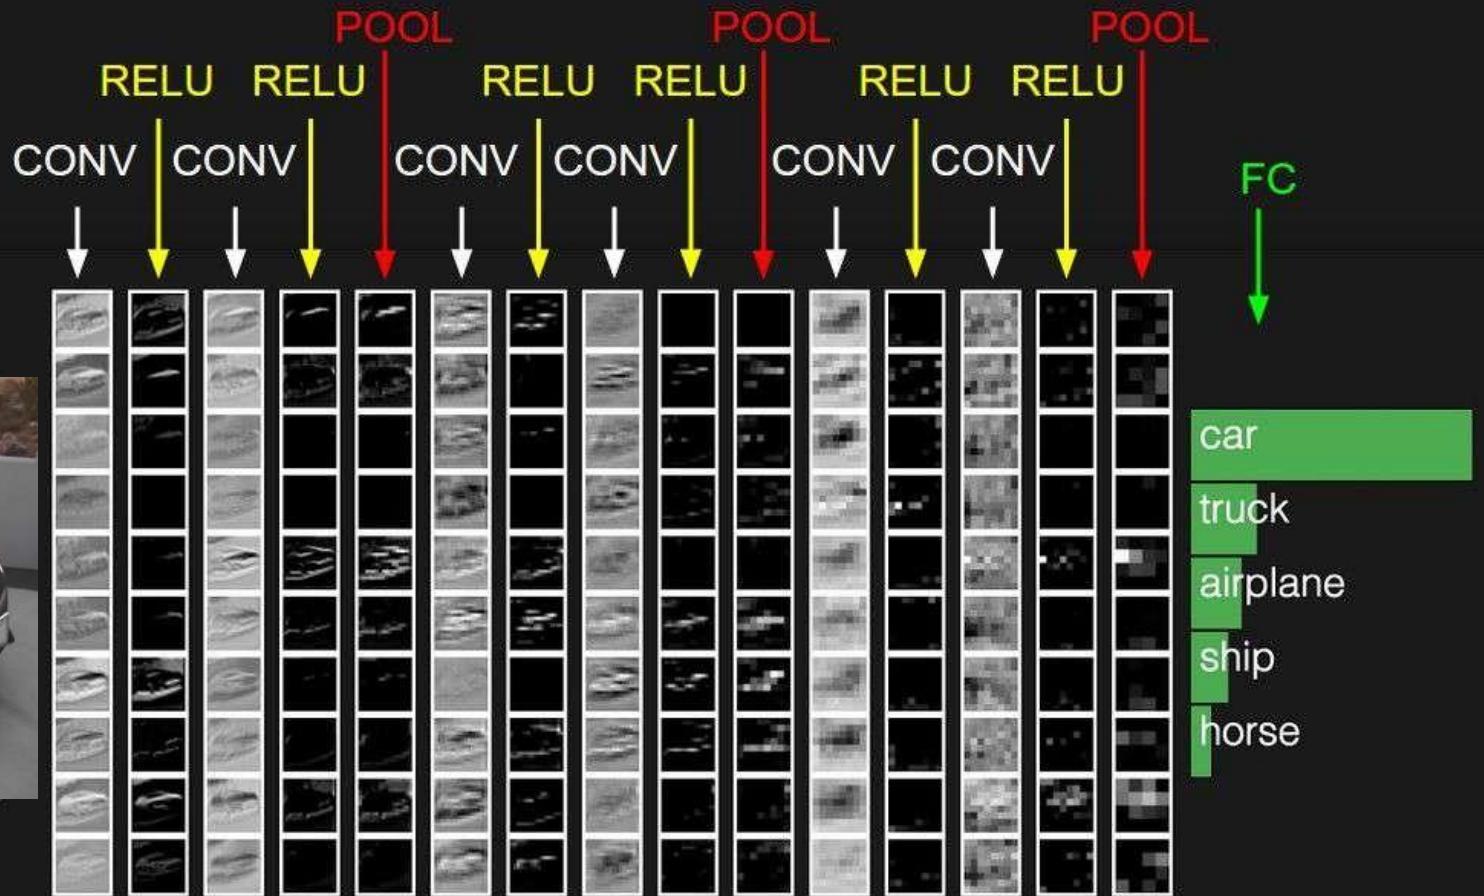
Podsetnik: Potpuno Povezani Sloj

32x32x3 slika -> razvijamo u 3072×1

Svaki neuron
je povezan sa
celom slikom.

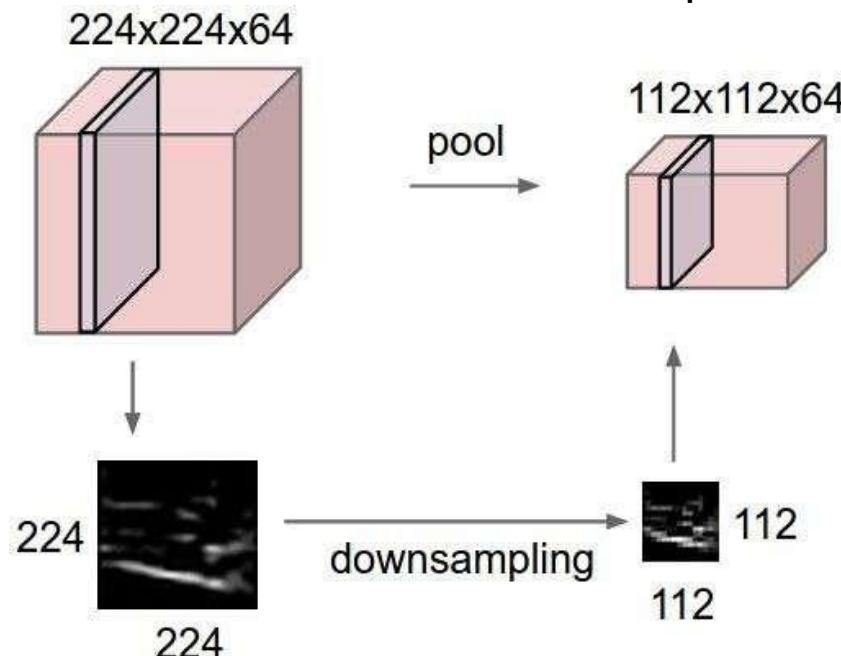


Nismo još objasnili sledeće slojeve: POOL/FC



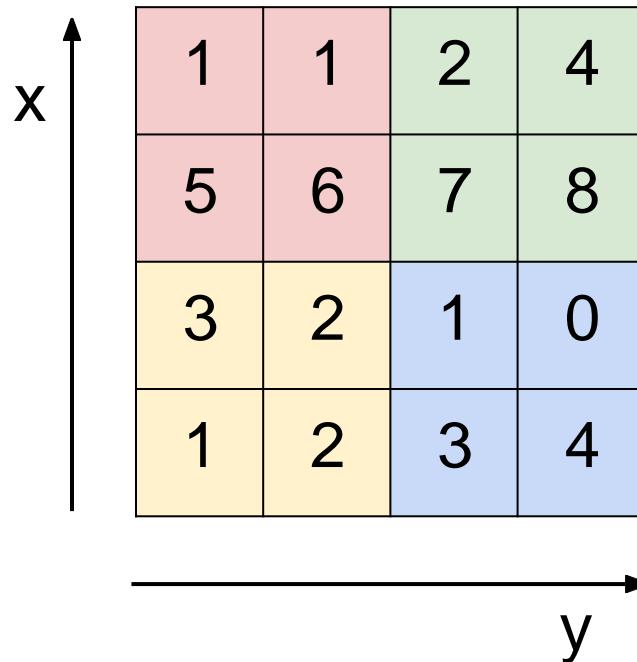
Pooling Sloj

- smanjujemo dimenzionalnost slojeva, sa ciljem da možemo lakše da obučimo mrežu (manje parametra), ali da pri tom sačuvamo što više informacija
- primenjuje se nezavisno na svaku aktivacionu mapu:



MAX POOLING

Jedna aktivaciona mapa dubine 1



max pool sa 2x2
fliterima i korakom 2

A horizontal arrow points from the input map to the output map. The output map is a 2x2 matrix. It has a pink cell containing 6 (top-left), a green cell containing 8 (top-right), a yellow cell containing 3 (bottom-left), and a blue cell containing 4 (bottom-right).

6	8
3	4

- Accepts a volume of size $W_1 \times H_1 \times D_1$
- Requires three hyperparameters:
 - their spatial extent F ,
 - the stride S ,
- Produces a volume of size $W_2 \times H_2 \times D_2$ where:
 - $W_2 = (W_1 - F)/S + 1$
 - $H_2 = (H_1 - F)/S + 1$
 - $D_2 = D_1$
- Introduces zero parameters since it computes a fixed function of the input
- Note that it is not common to use zero-padding for Pooling layers

Tipična podešavanja:

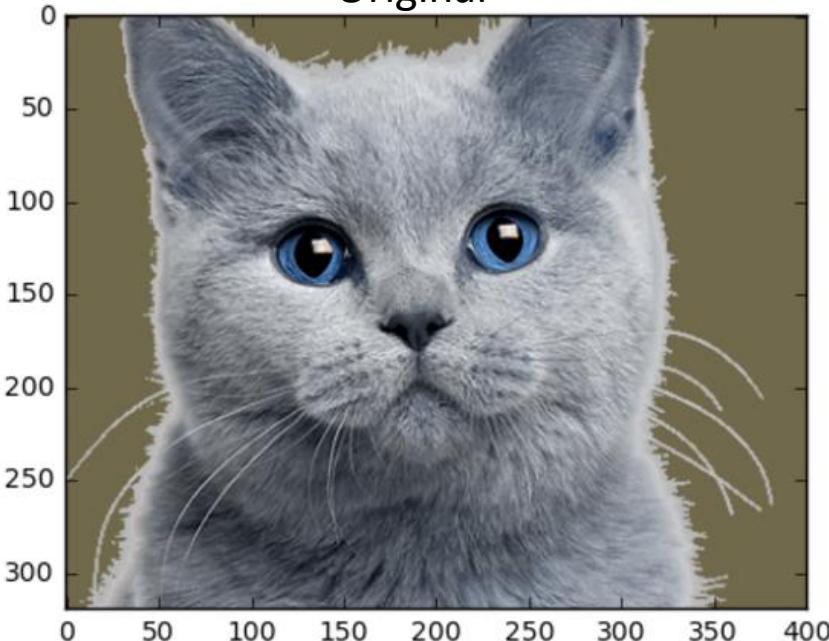
- Accepts a volume of size $W_1 \times H_1 \times D_1$
- Requires three hyperparameters:
 - their spatial extent F ,
 - the stride S ,
- Produces a volume of size $W_2 \times H_2 \times D_2$ where:
 - $W_2 = (W_1 - F)/S + 1$
 - $H_2 = (H_1 - F)/S + 1$
 - $D_2 = D_1$
- Introduces zero parameters since it computes a fixed function of the input
- Note that it is not common to use zero-padding for Pooling layers

$F = 2, S = 2$

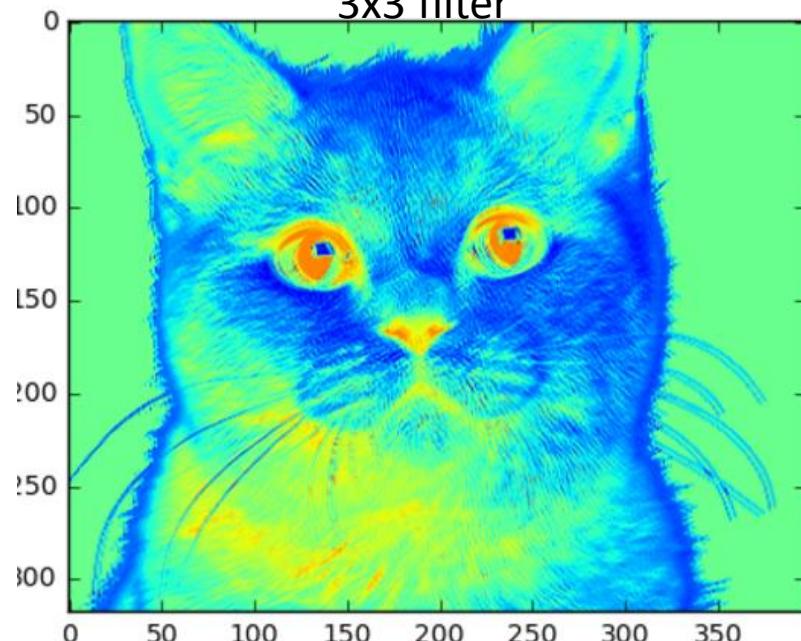
$F = 3, S = 2$

Primer jednog konvolucionog filtera 1

Original



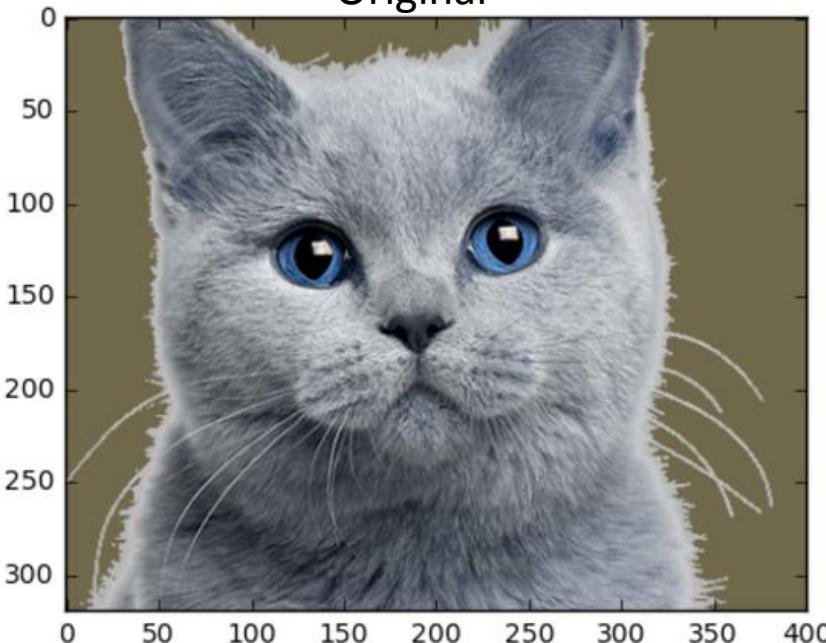
3x3 filter



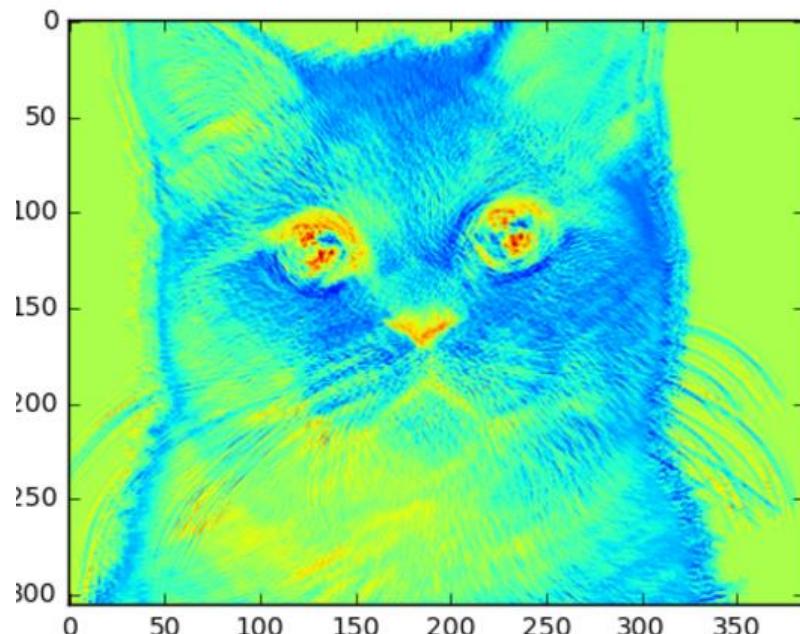
Izvor: <https://hackernoon.com/visualizing-parts-of-convolutional-neural-networks-using-keras-and-cats-5cc01b214e59>

Primer jednog konvolucionog filtera 2

Original

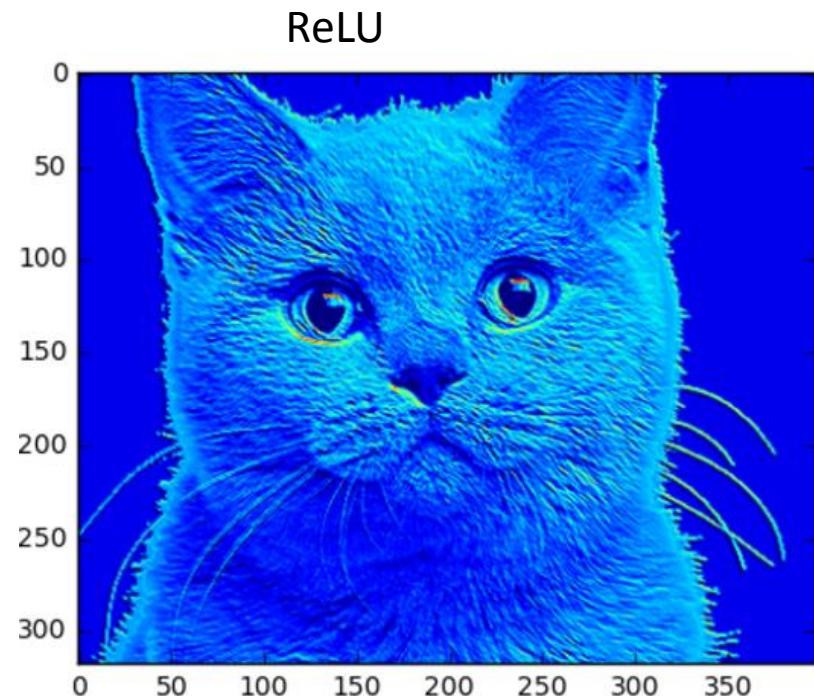
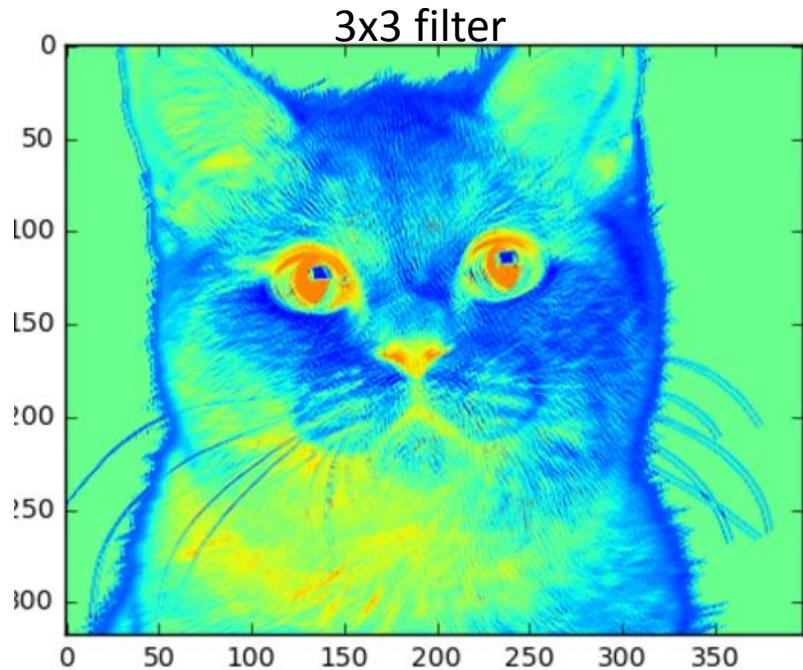


15x15 filter



Izvor: <https://hackernoon.com/visualizing-parts-of-convolutional-neural-networks-using-keras-and-cats-5cc01b214e59>

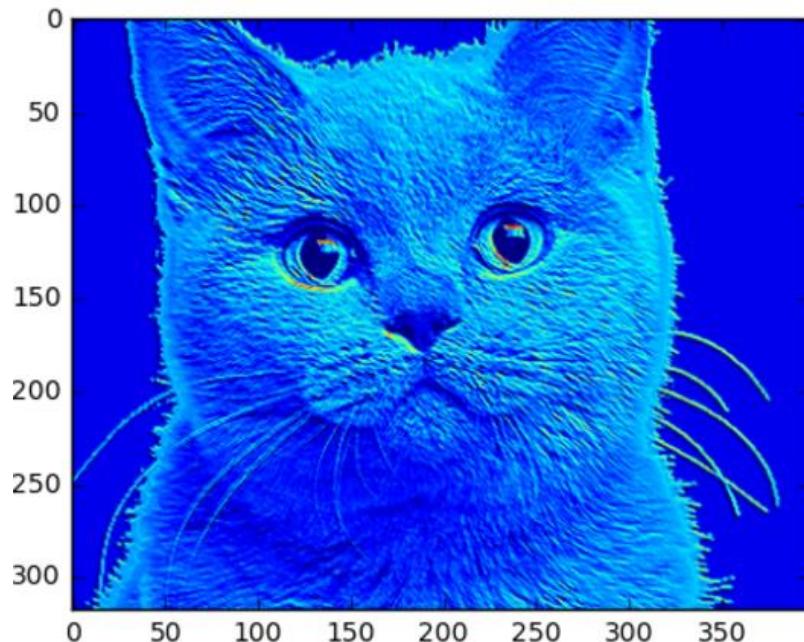
Primer jednog konvolucionog filtera 3, ReLU posle 3x3 filtera



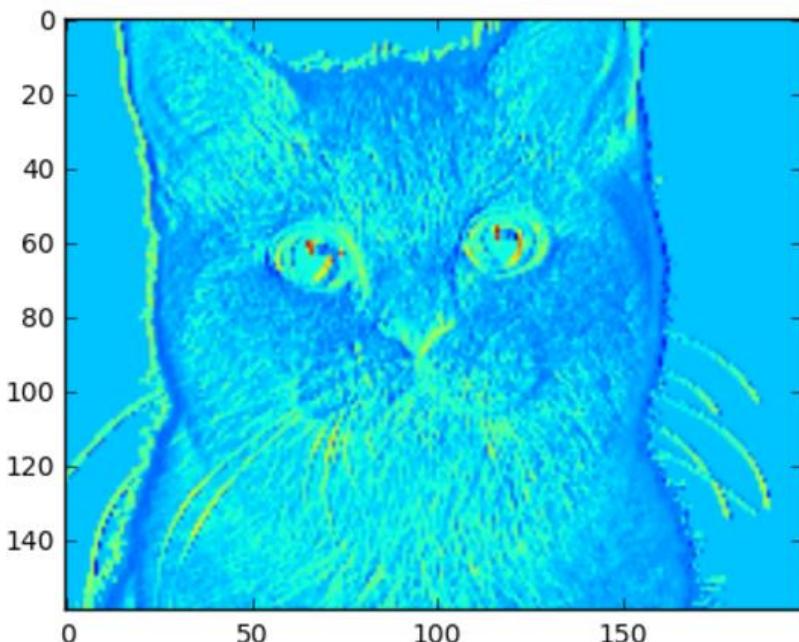
Izvor: <https://hackernoon.com/visualizing-parts-of-convolutional-neural-networks-using-keras-and-cats-5cc01b214e59>

Primer jednog konvolucionog filtera 4, Max Pooling 2x2 posle ReLU

ReLU



Max Pooling 2x2



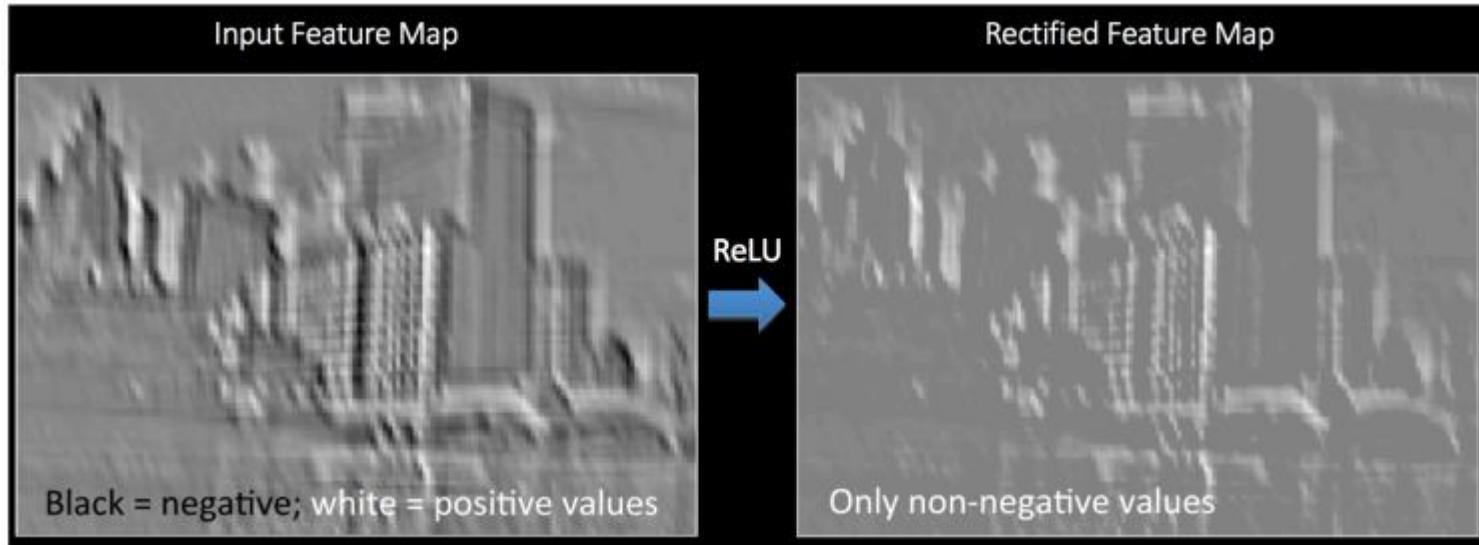
Izvor: <https://hackernoon.com/visualizing-parts-of-convolutional-neural-networks-using-keras-and-cats-5cc01b214e59>

Primer još jednog konvolucionog filtera 1



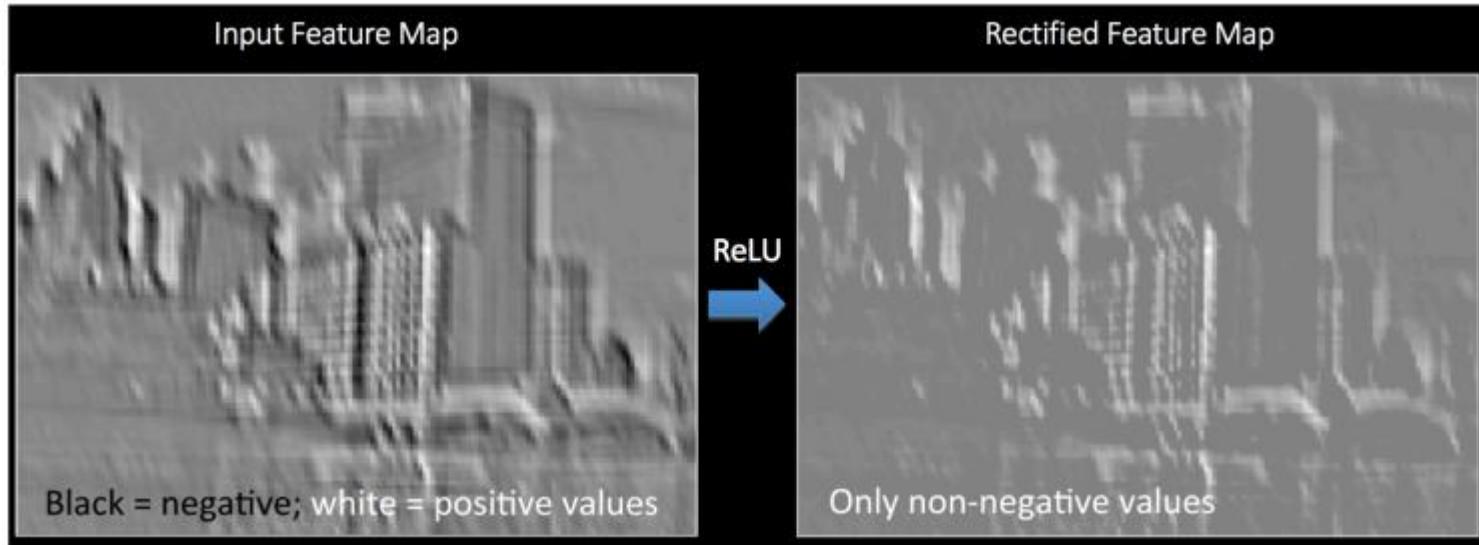
Izvor: <https://ujjwalkarn.me/2016/08/11/intuitive-explanation-convnets/>

Primer još jednog konvolucionog filtera 2, ReLU



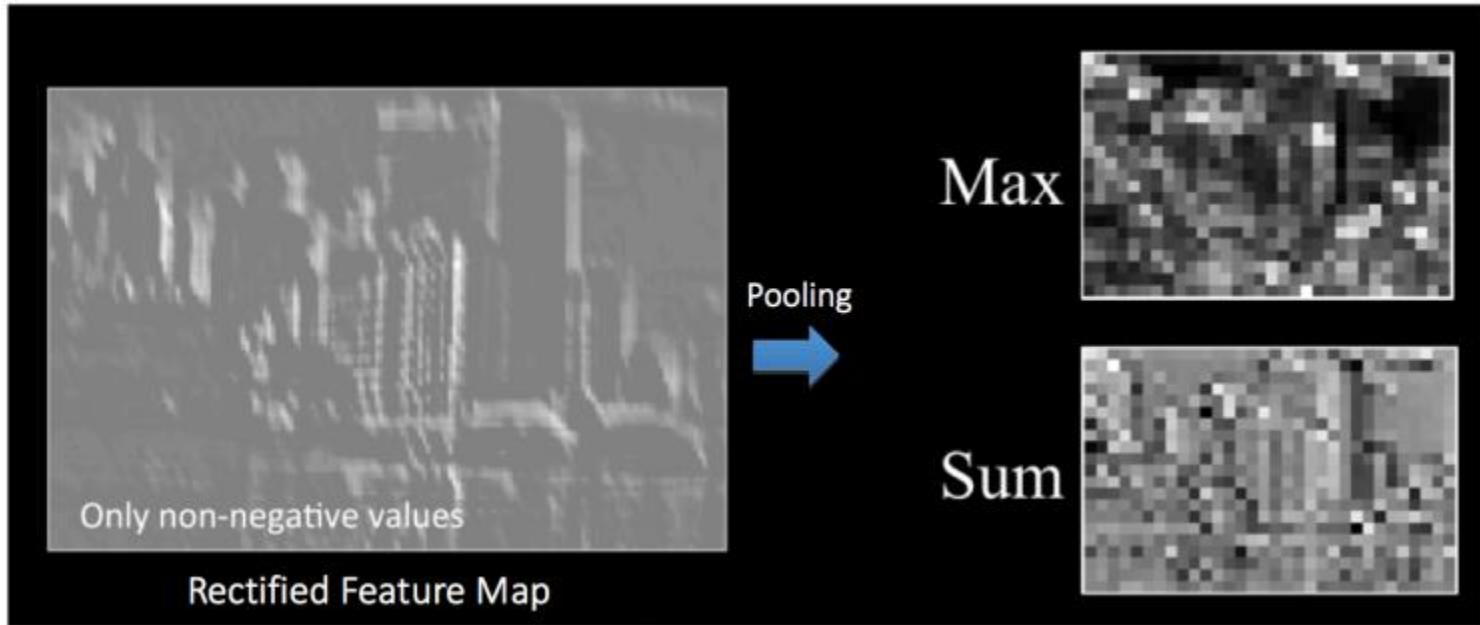
Izvor: <https://ujjwalkarn.me/2016/08/11/intuitive-explanation-convnets/>

Primer još jednog konvolucionog filtera 2, ReLU



Izvor: <https://ujjwalkarn.me/2016/08/11/intuitive-explanation-convnets/>

Primer još jednog konvolucionog filtera 3, Pooling

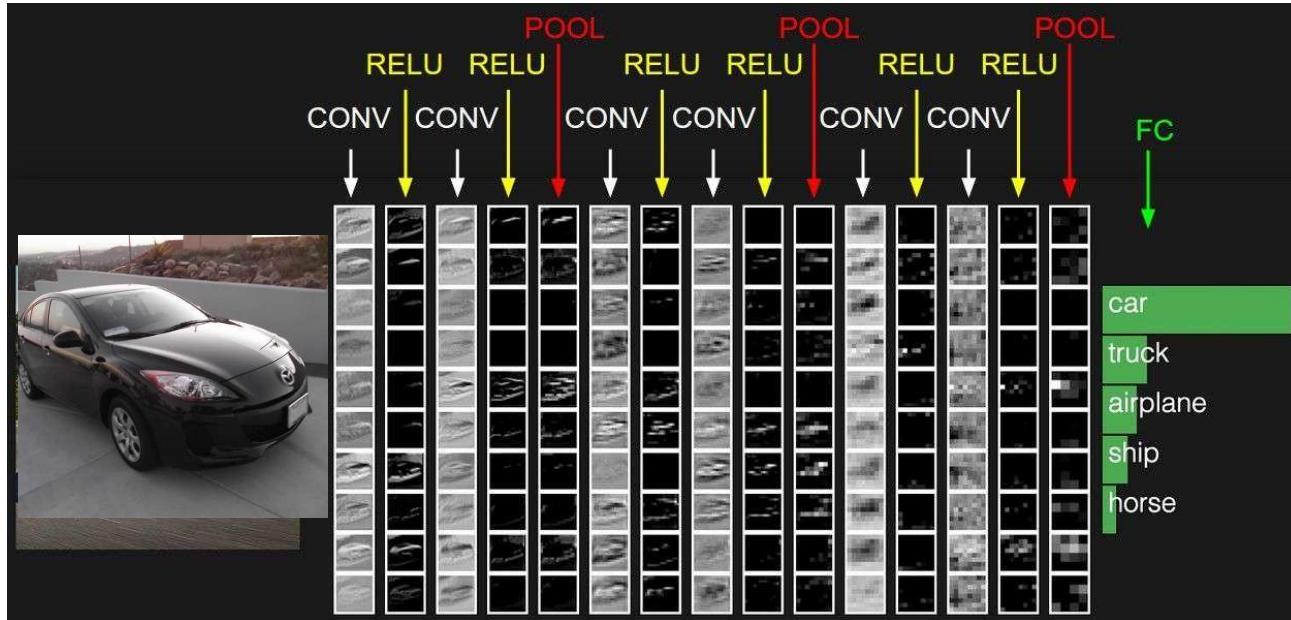


Izvor: <https://ujjwalkarn.me/2016/08/11/intuitive-explanation-convnets/>

Potpuno Povezani Sloj

Fully Connected Layer (FC layer)

- Sadrži neurone koji su potpuno povezani sa svojim ulazom, kao u tipičnoj neuronskoj mreži

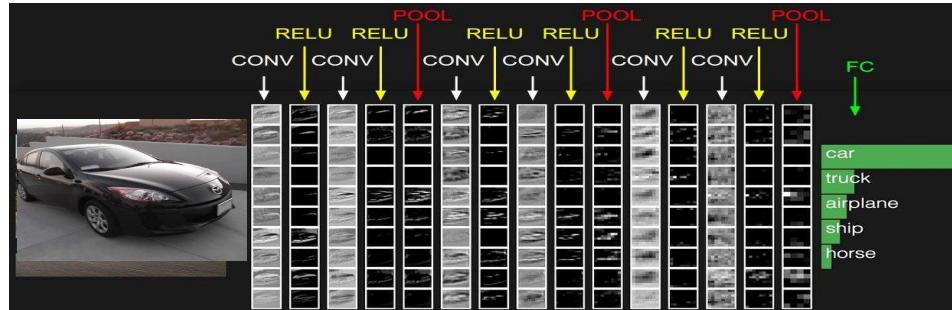


Potpuno Povezani Sloj

Fully Connected Layer (FC layer)

- Sadrži neurone koji su potpuno povezani sa svojim ulazom, kao u tipičnoj neuronskoj mreži
- Konvolutivne slojeve možemo da shvatimo kao ekstrakciju osobina iz slike koje pomoću kojih onda obučavamo potpuno povezanu neuronsku mrežu na kraju.

Zašto nisu svi slojevi potpuno povezani?



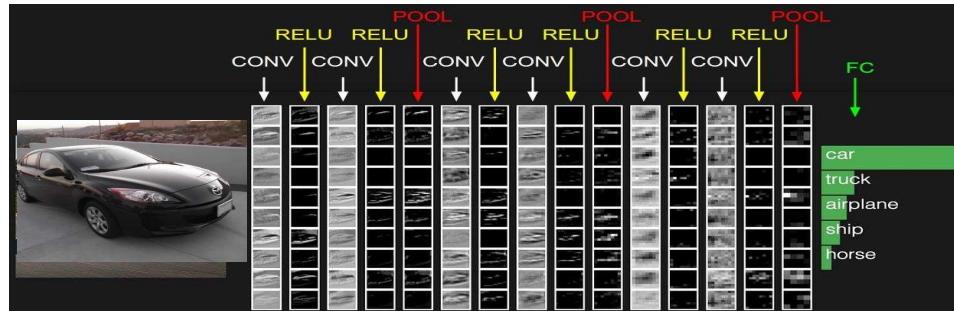
Potpuno Povezani Sloj

Fully Connected Layer (FC layer)

- Sadrži neurone koji su potpuno povezani sa svojim ulazom, kao u tipičnoj neuronskoj mreži
- Konvolutivne slojeve možemo da shvatimo kao ekstrakciju osobina iz slike koje pomoću kojih onda obučavamo potpuno povezanu neuronsku mrežu na kraju.

Zašto nisu svi slojevi potpuno povezani?

Takava mreža bi imala previše parametra da bi realno mogla da se obuči!



[ConvNetJS demo: obučavanje na CIFAR-10]

[ConvNetJS CIFAR-10 demo](#)

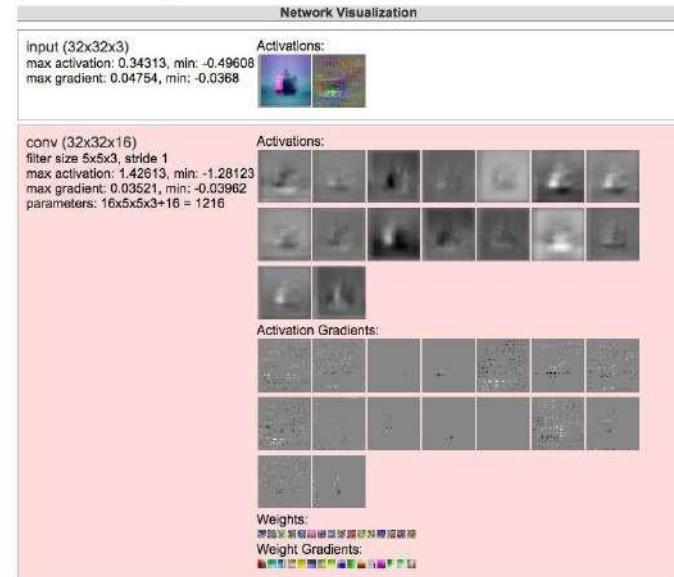
Description

This demo trains a Convolutional Neural Network on the [CIFAR-10 dataset](#) in your browser, with nothing but Javascript. The state of the art on this dataset is about 90% accuracy and human performance is at about 94% (not perfect as the dataset can be a bit ambiguous). I used [this python script](#) to parse the [original files](#) (python version) into batches of images that can be easily loaded into page DOM with img tags.

This dataset is more difficult and it takes longer to train a network. Data augmentation includes random flipping and random image shifts by up to 2px horizontally and vertically.

By default, in this demo we're using Adadelta which is one of per-parameter adaptive step size methods, so we don't have to worry about changing learning rates or momentum over time. However, I still included the text fields for changing these if you'd like to play around with SGD+Momentum trainer.

Report questions/bugs/suggestions to [@karpathy](#).



<http://cs.stanford.edu/people/karpathy/convnetjs/demo/cifar10.html>

Rezime

- Kovolutivne mreže slažu CONV,POOL i FC slojeve
- Trenutno je trend ka manjim filterima, a dubljim mrežama
- Eksperimentiše se i sa uklanjanjem POOL/FC slojeva (ostaju samo CONV)
- Tipična arhitektura izgleda ovako:

$[(CONV-RELU)^*N-POOL?]^*M-(FC-RELU)^*K, SOFTMAX$

gde je $N \sim 5$, M je veliko, $0 \leq K \leq 2$.

- ali neke od trenutno najboljih mreža ResNet/GoogLeNet menjaju ovu arhitekturu